

Text Preprocessing and Tokenization

PREPARED BY: AHMAD ALAA ALDINE

Text Normalization

LOWERCASING, PUNCTUATION REMOVAL, STEMMING VS.
LEMMATIZATION.

How many words in a sentence?

They picnicked by the pool, then they lay back
on the grass and looked at the stars.

17 words

- if we don't count punctuation marks as words

19 if we count punctuation

Vocabulary, Type and Instance

They picnicked by the pool, then they lay back
on the grass and looked at the stars.

Vocabulary: all unique words in a given text or corpus

Type: an element of the vocabulary V

- The number of types is the vocabulary size $|V|$

Instance: an instance of that type in a given text.

- 15 types and 17 instances (if we ignore punctuation).

More questions: Are **They** and **they** the same type?

The Problem: A Vocabulary Explosion

In Natural Language Processing, treating every word variation as a unique token creates massive, sparse vocabularies.

As text corpora grow, the number of unique word types grows without a clear bound, posing a huge problem for any computational model.

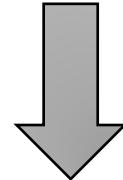
Key Takeaway: To build robust models, we must group different forms of the same word. This process is called **Text Normalization**.

Lowercasing

Lowercasing in NLP is a text preprocessing step that converts all text to lowercase to standardize words (e.g., "They" becomes "they")

They picnicked by the pool, then they lay back
on the grass and looked at the stars.

Lowercasing



they picnicked by the pool, then they lay back
on the grass and looked at the stars.

$|V| = 15$
 $N = 17$

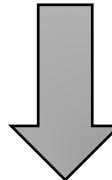
$|V| = 14$
 $N = 17$

- Reduce Vocabulary Size

Proper names may lose identity!

Apple released a new product in the US

Lowercasing



apple released a new product in the us



Model can no longer distinguish:

- **Apple (Company) vs apple (fruit)**
- **US (Country) vs us (pronoun)**

Punctuation Removal

Punctuation removal in NLP is a crucial text preprocessing step that strips characters like periods, commas, and question marks.

We need to carefully choose the list of punctuation which we are going to discard based on the use case.

For example, Python's **string** module contains the following list of punctuation.

```
>>> import string  
>>> string.punctuation  
' !"#$%&\' () *+, -./ :; <=>?@ [ \\ ] ^ ` { | } ~'
```

You can also add other special symbols to list which you want to discard.

Why Remove Punctuation in Text Preprocessing?

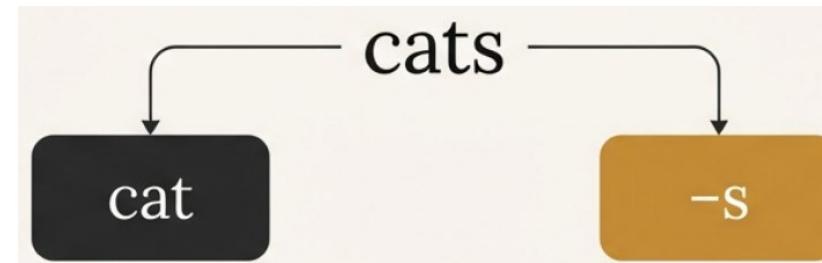
Reduces Noise: Punctuation adds characters that often lack meaning.

Simplifies Data: Creates fewer unique types, reducing vocabulary size and computational load.

Improves Accuracy: Prevents models from learning false distinctions (e.g., "hello" vs. "hello!").

The Building Blocks of Meaning: Introducing Morphemes

A **morpheme** is the minimal meaning-bearing unit in a language. Words are composed of one or more morphemes.



(the **root** morpheme,
supplying the main meaning)

(an **affix**, indicating plural)

Morphemes are everywhere. The sentence below is segmented to show its morphemic structure:

Doc work - ed care - ful - ly wash - ing the glass - es

Two Classes of Affixes

Inflectional Morphemes

Function: Create different grammatical *forms of the same* word. They play a syntactic role (e.g., tense, number).

Examples:

-s / -es for plural nouns (cat cats)

-ed for past tense verbs (work worked)

Derivational Morphemes

Function: Create *new* words, often of a different grammatical class.

Examples:

-ly to create an adverb (careful carefully)

-ation to create a noun (organize organization)

Another Type of Morpheme: Clitics

A morpheme that acts syntactically like a word but:

- is reduced in form
- and attached to another word

English: 've in I've ('ve can't appear alone)

English: 's in the teacher's book

French: l' in l'opera

Arabic: ب 'by/with', و 'and'.

Stemming and Lemmatization

Lemmatization and Stemming are normalization techniques used in NLP to reduce the word to its' base (root) form.

Stemming uses a heuristic approach to reduce the word to its' root form.

Lemmatization uses linguistic knowledge (dictionaries, context) to find the actual dictionary word (lemma).

	Stemming	Lemmatization
Emailing	email	email
Replies	repli	reply
Reporting	report	reporting
Presentations	present	presentation
Meeting	meet	meeting
Scheduling	schedul	schedule

Stemming vs. Lemmatization

Stemming

Method: Heuristic, rule-based suffix removal (e.g., Porter Stemmer).

Output: A "root" that might not be a real word (e.g., "caring" -> "car", "studies" -> "studi").

Pros: Fast, efficient for large datasets, lower computational cost.

Cons: Lower accuracy, can lose semantic meaning.

Lemmatization

Method: Uses dictionaries (lexicon) and Part-of-Speech (POS) tagging to find the canonical form (lemma) (e.g., WordNet Lemmatizer).

Output: A valid, meaningful dictionary word (e.g., "caring" -> "care", "better" -> "good").

Pros: High accuracy, preserves meaning, consistent output.

Cons: Slower, computationally intensive, requires more resources (dictionaries).

POS tagging: assigning a grammatical category to each word.

To be covered later →

The Core Trade-Off: Speed vs. Accuracy

Stemming: FAST & SIMPLE

Reduces feature space quickly.

Output is not linguistically correct or easily interpretable.

Lemmatization: ACCURATE & INTERPRETABLE

Produces valid words, enabling deeper semantic analysis.

Computationally expensive and requires more linguistic resources (lexicon, POS tags).

Your choice depends entirely on the requirements of your specific NLP application.

Tokenization Strategies

A JOURNEY FROM RULE-BASED METHODS TO THE SUBWORD REVOLUTION

Text is unstructured Machine understand numbers

The first step in any modern NLP pipeline is converting raw text into a sequence of discrete units, called tokens.

This fundamental task is called tokenization.

It is the bridge between human language and machine understanding.



Treat Words as Tokens: Early methods

Early methods are rule-based methods that rely on handcrafted rules to define word boundaries.

The **Penn Treebank** tokenization standard, for example, uses a sophisticated set of rules to handle punctuation and clitic contractions in English.



The Penn Treebank

The **Penn Treebank tokenizer** uses regular expressions to tokenize text.

This tokenizer performs the following steps:

- split standard contractions, e.g. ``don't`` -> ``do n't`` and ``they'll`` -> ``they 'll``
- treat most punctuation characters as separate tokens
- split off commas and single quotes, when followed by whitespace
- separate periods that appear at the end of line

Examples of **Penn Treebank** regular expressions to handle clitics:

```
ENDING_QUOTES = [
    (re.compile(r'''')), " '' "),
    (re.compile(r''''), " '' "),
    (re.compile(r"([^\ ])('['sS]|\ '['mM]|\ '['dD]|\ ') "), r"\1 \2 "),
    (re.compile(r"([^\ ])('ll| 'LL| 're| 'RE| 've| 'VE| 'n't| 'N'T) "), r"\1 \2 "),
]
```

Beyond Rule-Based Tokenization

RULE-BASED TOKENIZATION

Penn Treebank Tokenization

- Rule-based
- Deterministic
- Language-specific

Limitations

- Hard-coded rules
- Difficult to adapt to new languages
- Edge cases remain

STATISTICAL-BASED TOKENIZATION

- Learns token boundaries from data
- Uses context and frequency
- Adapts better across domains

Examples

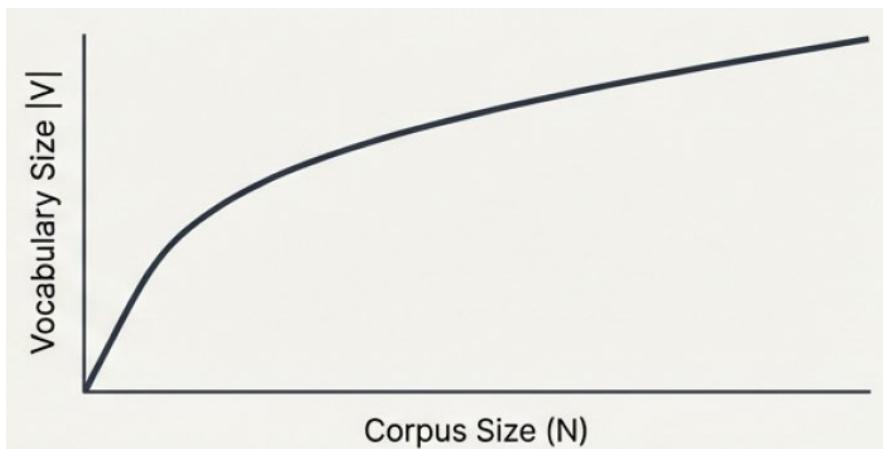
- Learned tokenizers in modern NLP tools (**BPE**)
- Punkt tokenizer (sentence-level)
- HMM / CRF-based word segmentation

Word-Level Tokenization Limitations

Relying on words as tokens presents two critical failures that limit **scalability** and **cross-lingual applicability**.

Scalability

The vocabulary explosion. The number of unique words in a language grows unboundedly as we process more text.



Cross-lingual Applicability

How do you segment languages without spaces?

In Chinese, 姚明进入总决赛 “Yao Ming reaches the finals” has multiple valid word tokenization.

姚明进入总决赛						
姚明	进入	总	决赛 (3 tokens)			
姚	明	进	入	总	决	赛 (5 tokens)
姚	明	进	入	总	决	赛 (7 tokens)

[UNK] – Unknown Words

Even a massive training corpus will not contain every word.

UNK represents a word that is not in the model's vocabulary, also known as out-of-vocabulary (oov) words.

The UNK token serves as a universal placeholder for all such words, allowing the model to process any input text without encountering errors.

Training Vocabulary = {i, love, nlp}

New Input Sentence: "I love transformers"

After Tokenization: ["i", "love", "UNK"]

"transformers" was never seen during training

Problem caused by [UNK]

- Loss of semantic information
- Different words become indistinguishable

Solution: Subword Tokenization

Subword tokenization: represent words with sequences of sub-units.

It addresses the failures of word-based methods by creating a fixed-size vocabulary of reusable parts.

These subwords are often morphemes (meaningful units like '-er', 're-') or other frequently occurring character sequences.

We're using a transformer-based model.

We're using a transform ##er [-] based model

Subword Tokenization: Advantages

Eliminates '[UNK]'

- Any new or rare word can be constructed from known subword pieces

Fixed Vocabulary

- The vocabulary size is a fixed hyperparameter, preventing uncontrolled growth.

Captures Morphology

- Related words like 'low' 'lower' and 'lowest' can share the 'low' subword token, allowing the model to see their connection.

Modern Tokenizers

Byte-Pair Encoding (BPE)

- Iteratively merges the most frequent pairs of characters or character sequences into a single token.
- It is the standard for models like **GPT** and **LLaMA**.

WordPiece

- Similar to BPE but merges pairs based on maximizing the likelihood of the training data rather than just frequency.
- It is used by **BERT** and **DistilBERT**.

Unigram Language Modeling (ULM)

- A probabilistic approach that starts with a large vocabulary and removes tokens that least impact the overall likelihood of the training set.

Byte-Pair Encoding (BPE)

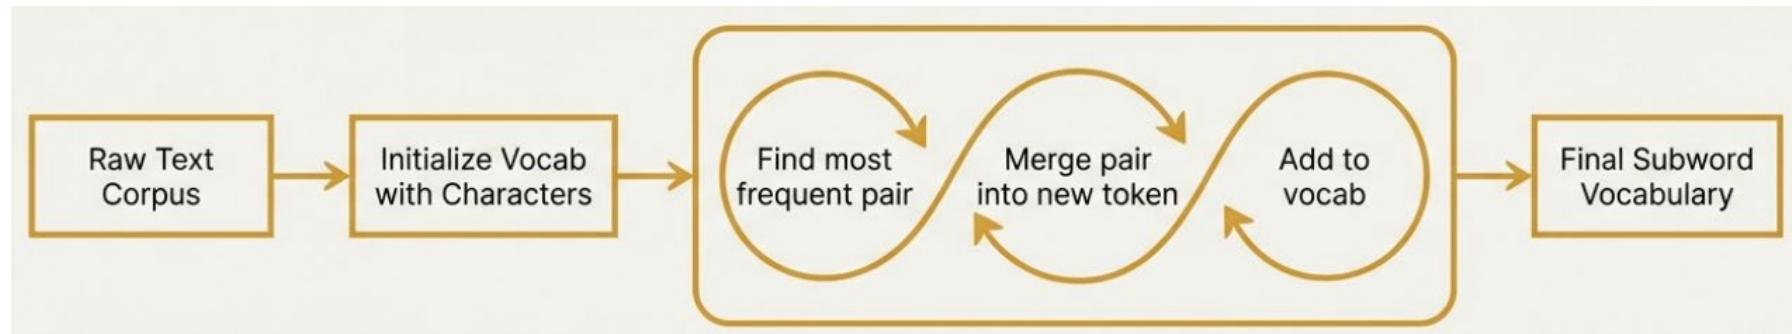
BPE builds a vocabulary by learning the most efficient "text compression".

Originally a data compression algorithm, it was adapted for NLP tokenization.

Its core principle is powerful and requires no linguistic expertise.

Start with individual characters, and iteratively merge the most frequently occurring adjacent pair of tokens.

The process is repeated for a predetermined number of merges, building up a vocabulary of common subwords.



BPE Training Algorithm

```
function BYTE-PAIR ENCODING(strings C, number of merges k) returns vocab V
    V all unique characters in C                                # initial set of tokens is characters
    for i = 1 to k do                                         # merge tokens k times
         $t_l, t_r \leftarrow$  Most frequent pair of adjacent tokens in C
         $t_{NEW} \leftarrow t_l + t_r$                                # make new token by concatenating
         $V \leftarrow V + t_{NEW}$                                  # update the vocabulary
        Replace each occurrence of  $t_l, t_r$  in C with  $t_{NEW}$       # and update the corpus
    return V
```

Training the Vocabulary from Data (Part 1)

Let's trace the BPE training process on a small, synthetic corpus.

Corpus: “set new new renew reset renew”

First, break up the corpus into words, with leading whitespace, together with their counts.

Merges are only allowed within the initial word boundaries.

Initial State:

corpus	vocabulary
2 « n e w	«, e, n, r, s, t, w
2 « r e n e w	
1 s e t	
1 « r e s e t	

Training the Vocabulary from Data (Part 2)

Step 1: First merge

The BPE training algorithm first counts all pairs of adjacent symbols:

the most frequent is the pair **n e** because it occurs in “**new**” (frequency of 2) and “**renew**” (frequency of 2) for a total of 4 occurrences.

We then merge these symbols, treating **ne** as one symbol, and count again:

corpus	vocabulary
2 ↳ n e w	↳, e, n, r, s, t, w, ne
2 ↳ r e n e w	
1 s e t	
1 ↳ r e s e t	

Training the Vocabulary from Data (Part 3)

Step 2: Second merge

Now the most frequent pair is ne w (total count=4), which we merge.

corpus	vocabulary
2 _ new	_ , e, n, r, s, t, w, ne, new
2 _ r e new	
1 s e t	
1 _ r e s e t	

Step 3: Third merge

Next _ and r (total count of 3) get merged to _r.

corpus	vocabulary
2 _ new	_ , e, n, r, s, t, w, ne, new, _r
2 _re new	
1 s e t	
1 _re s e t	

Training the Vocabulary from Data (Part 4)

Step 4: Fourth merge

Next „r and e (total count of 3) get merged to „re.

corpus	vocabulary
2 „ new	„, e, n, r, s, t, w, ne, new, „r, „re
2 „re new	
1 s e t	
1 „re s e t	

If we continue, the next merges are:

merge	current vocabulary
(„, new)	„, e, n, r, s, t, w, ne, new, „r, „re, „new
(„re, new)	„, e, n, r, s, t, w, ne, new, „r, „re, „new, „renew
(s, e)	„, e, n, r, s, t, w, ne, new, „r, „re, „new, „renew, se
(se, t)	„, e, n, r, s, t, w, ne, new, „r, „re, „new, „renew, se, set

BPE Encoder

Once we've learned our vocabulary, the BPE encoder is used to tokenize a test sentence.

The encoder just runs on the test data the merges we have learned from the training data.

It runs them in the order we learned them.

Example:

Input word “newest”, Merge List: `_, e, n, r, s, t, w, ne, new, _r, _re, _new, _renew, se, set`

Initial representation: Each character is a token.

Tokens: “n” “e” “w” “e” “s” “t”

Step 1: Apply Merge Rule #1 (`n + e -> ne`)

Tokens: “ne” “w” “e” “s” “t”



Step 2: Apply Merge Rule #2 (`ne + w -> new`)

Tokens: “new” “e” “s” “t”

No more merge rules to apply

Sentence Segmentation

From Tokens to Sentences

Tokenization is not only about words.

Many NLP tasks operate at the sentence level such as **Translation**.

Sentence segmentation is the task of identifying sentence boundaries in text.

Dr. Smith arrived at 5 p.m. He left early.

Naive split fails

Intelligent segmentation required

Why Sentence Segmentation is Hard

Sentence segmentation depends on the language and the genre.

The most useful cues for segmenting a text into sentences in English written text tend to be punctuation, like periods “.”, question marks “?”, and exclamation points “!”.

Question marks “?” and exclamation points “!” are relatively unambiguous markers of sentence boundaries, and simple rules can segment sentences when they appear.

Period “.” is quite **ambiguous**

- Sentence boundary
- Abbreviations like Inc. or Dr.
- Numbers like .02% or 4.3

Sentence Segmentation Methods

Many English sentence segmentation methods work by first deciding whether a period is part of the word or is a sentence-boundary marker.

- Often based on deterministic rules ([Stanford CoreNLP sentence splitter](#))
- But sometimes via machine learning ([NLTK Punkt tokenizer](#))

Key Takeaway: Rule-based methods are fast and simple, but **they are language-specific**.
Machine Learning methods generalize better.

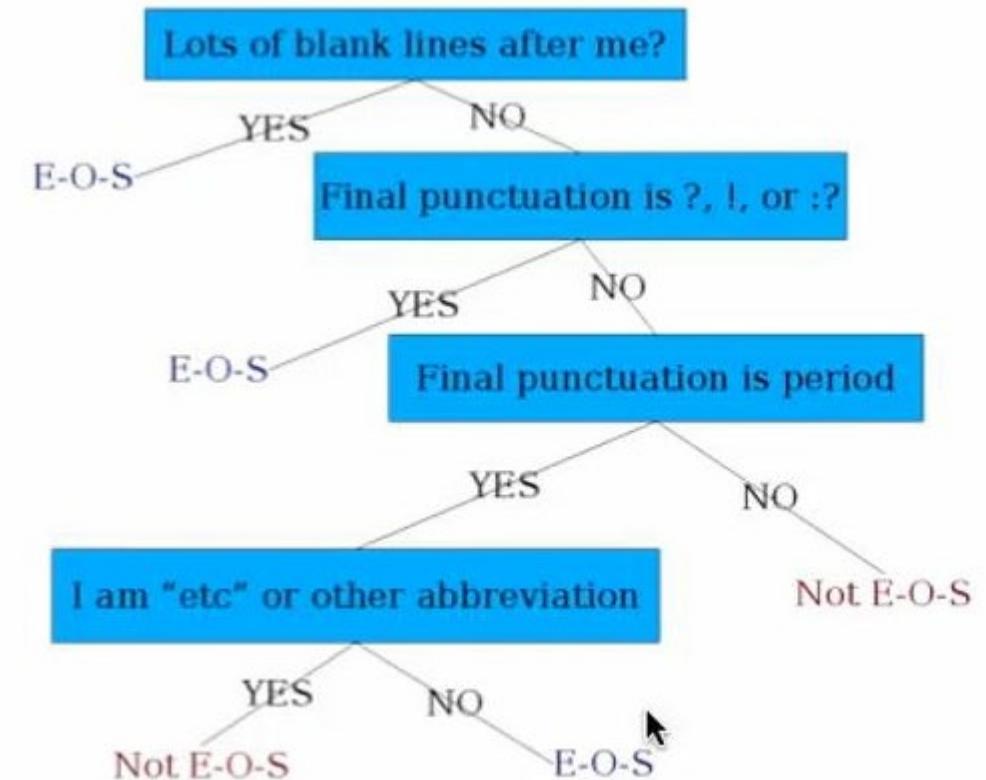
The Simplest Method: Decision Tree

A Decision Tree is simply an if-else statement.

The interesting research in a decision tree is choosing the features.

Lots of modifications can be made to make a more sophisticated decision tree such as:

- Case of word with “.”: Upper, Lower, Cap, Number
- Case of word after “.”: Upper, Lower, Cap, Number
- Length of word with “.”



Stanford CoreNLP Sentence Splitter: Core Rules

A sentence *may* end at: “.”, “!” and “?”

Periods do NOT end a sentence if they belong to:

- Titles
- Common Abbreviations
- Acronyms

Periods inside numbers are ignored.

Periods after single-letter initials do not trigger a split.

- e.g: John F. Kennedy was elected.

Ignore sentence boundaries inside quotes and parentheses

- e.g: She said, "Stop."

Punkt Sentence Tokenizer

Punkt is an **unsupervised** statistical sentence segmentation algorithm.

- Learns abbreviations and sentence boundaries
- No manually labeled data required

Punkt decides based on context and frequency however a period “.” is:

- Sentence end
- Abbreviation
- Decimal point

Training Phase & Decision Phase

What Punkt Learns Automatically

- Common abbreviations (Dr., Mr., etc.)
- Typical sentence starters
- Token context around punctuation

Decision Rules

- Is the token frequently abbreviated?
- Does the next token look like a sentence start?
- Is capitalization consistent?

**Punkt assigns probability,
not hard rules**

Strengths and Limitations of Punkt

STRENGTHS

Language-adaptive

No labeled data

Handles abbreviations well

LIMITATIONS

Needs sufficient training text

Domain-specific abbreviations may fail

Not designed for informal text (tweets, chats)