

Sequence Labeling: POS Tagging and NER

PREPARED BY: AHMAD ALAA ALDINE

POS Tagging

Parts Of Speech – History

The idea that words can be classified into grammatical categories return back to 100 B.C. (**long time ago**)

Eight parts of speech attributed to Dionysius Thrax of Alexandria:

- noun, verb, pronoun, preposition, adverb, conjunction, participle, article
- These categories are still relevant for NLP today.

Two classes of words: Open vs. Closed

Closed class words

- Relatively fixed membership
- Usually **function** words: short, frequent words with grammatical function
 - determiners: *a, an, the*
 - pronouns: *she, he, I*
 - prepositions: *on, under, over, near, by, ...*

Open class words

- Usually **content** words: Nouns, Verbs, Adjectives, Adverbs
 - Plus interjections: *oh, ouch, uh-huh, yes, hello*
- New nouns and verbs like *iPhone* or *to fax*

Open class ("content") words

Nouns

Proper

*Janet
Italy*

Common

*cat, cats
mango*

Verbs

Main

*eat
went*

Adjectives

old green tasty

Adverbs

slowly yesterday

Numbers

*122,312
one*

Interjections

*Ow hello
... more*

Closed class ("function")

Determiners

the some

Conjunctions

and or

Pronouns

they its

Auxiliary

*can
had*

Prepositions

to with

Particles

off up

... more

What is POS Tagging?

Part-of-Speech (POS) tagging is the process of **assigning grammatical categories** (like noun, verb, adjective) to each word in a sentence based on both its **definition and context**.

Example

The

quick

brown

fox

jumps

DET

ADJ

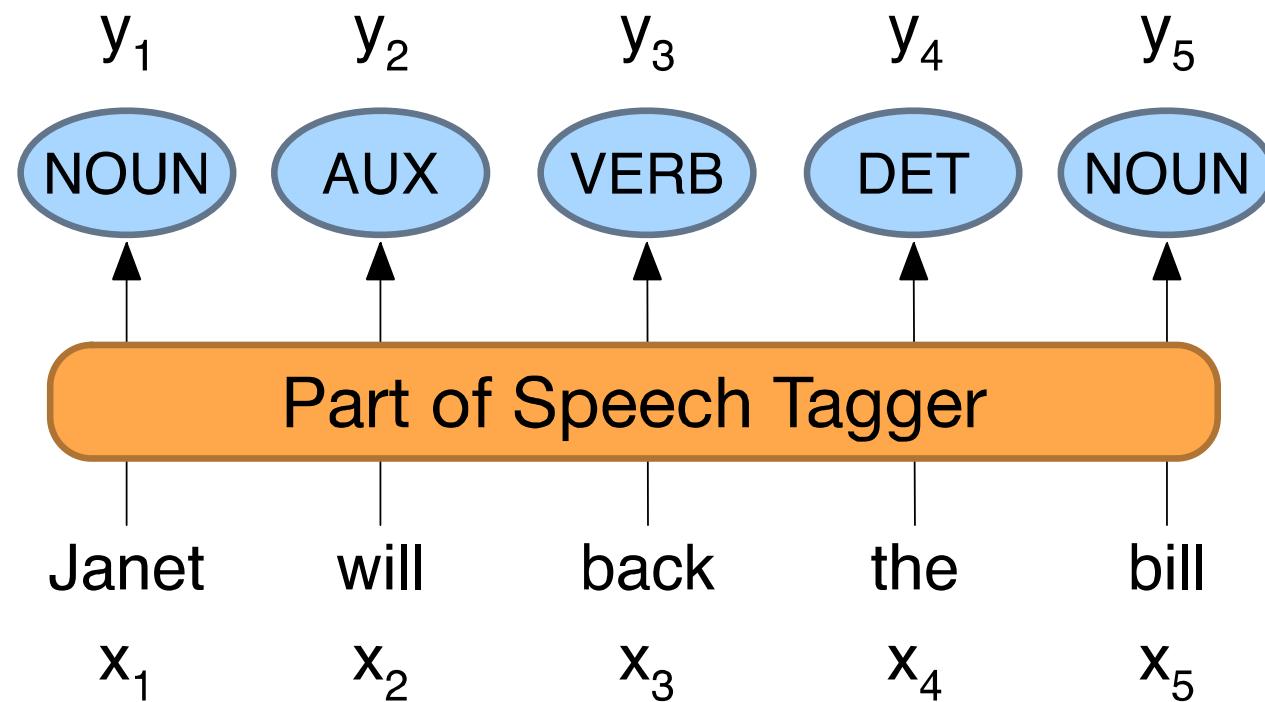
ADJ

NOUN

VERB

POS tagging – Sequence Labeling

Map from sequence x_1, \dots, x_n of words to y_1, \dots, y_n of POS tags



Common POS Tag Sets

Penn Treebank

45 tags

Most widely used (NN, VBZ, JJ, DT, etc.)

Universal Dependencies

17 tags

Language Independent (NOUN, VERB, ADJ, DET)

Example Penn Treebank Tags

- NN (noun, singular)
- VBZ (verb, 3rd person)
- JJ (adjective)
- DT (determiner)
- RB (adverb)
- PRP (pronoun)

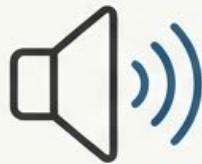
Universal Dependencies Tagset

Tag	Description	Example
Open Class	ADJ Adjective: noun modifiers describing properties	<i>red, young, awesome</i>
	ADV Adverb: verb modifiers of time, place, manner	<i>very, slowly, home, yesterday</i>
	NOUN words for persons, places, things, etc.	<i>algorithm, cat, mango, beauty</i>
	VERB words for actions and processes	<i>draw, provide, go</i>
	PROPN Proper noun: name of a person, organization, place, etc..	<i>Regina, IBM, Colorado</i>
	INTJ Interjection: exclamation, greeting, yes/no response, etc.	<i>oh, um, yes, hello</i>
Closed Class Words	ADP Adposition (Preposition/Postposition): marks a noun's spacial, temporal, or other relation	<i>in, on, by under</i>
	AUX Auxiliary: helping verb marking tense, aspect, mood, etc.,	<i>can, may, should, are</i>
	CCONJ Coordinating Conjunction: joins two phrases/clauses	<i>and, or, but</i>
	DET Determiner: marks noun phrase properties	<i>a, an, the, this</i>
	NUM Numeral	<i>one, two, first, second</i>
	PART Particle: a preposition-like form used together with a verb	<i>up, down, on, off, in, out, at, by</i>
	PRON Pronoun: a shorthand for referring to an entity or event	<i>she, who, I, others</i>
	SCONJ Subordinating Conjunction: joins a main clause with a subordinate clause such as a sentential complement	<i>that, which</i>
Other	PUNCT Punctuation	<i>; , ()</i>
	SYM Symbols like \$ or emoji	<i>\$, %</i>
	X Other	<i>asdf, qwfg</i>

Penn Treebank core 36 part-of-speech tags

Tag	Description	Example	Tag	Description	Example	Tag	Description	Example
CC	coord. conj.	<i>and, but, or</i>	NNP	proper noun, sing.	<i>IBM</i>	TO	infinitive to	<i>to</i>
CD	cardinal number	<i>one, two</i>	NNPS	proper noun, plu.	<i>Carolinas</i>	UH	interjection	<i>ah, oops</i>
DT	determiner	<i>a, the</i>	NNS	noun, plural	<i>llamas</i>	VB	verb base	<i>eat</i>
EX	existential ‘there’	<i>there</i>	PDT	predeterminer	<i>all, both</i>	VBD	verb past tense	<i>ate</i>
FW	foreign word	<i>mea culpa</i>	POS	possessive ending	<i>'s</i>	VBG	verb gerund	<i>eating</i>
IN	preposition/ subordin-conj	<i>of, in, by</i>	PRP	personal pronoun	<i>I, you, he</i>	VBN	verb past participle	<i>eaten</i>
JJ	adjective	<i>yellow</i>	PRP\$	possess. pronoun	<i>your</i>	VBP	verb non-3sg-pr	<i>eat</i>
JJR	comparative adj	<i>bigger</i>	RB	adverb	<i>quickly</i>	VBZ	verb 3sg pres	<i>eats</i>
JJS	superlative adj	<i>wildest</i>	RBR	comparative adv	<i>faster</i>	WDT	wh-determ.	<i>which, that</i>
LS	list item marker	<i>1, 2, One</i>	RBS	superlatv. adv	<i>fastest</i>	WP	wh-pronoun	<i>what, who</i>
MD	modal	<i>can, should</i>	RP	particle	<i>up, off</i>	WP\$	wh-possess.	<i>whose</i>
NN	sing or mass noun	<i>llama</i>	SYM	symbol	<i>+, %, &</i>	WRB	wh-adverb	<i>how, where</i>

Why Part of Speech Tagging Matters



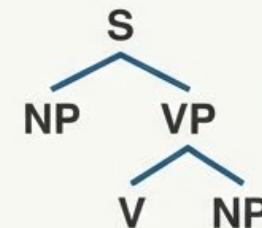
Text-to-Speech (TTS)

Pronunciation depends on the tag.
“OB-ject” (**Noun**) vs “ob-JECT”
(**Verb**).
“CON-tent” (**Noun**) vs “con-TENT”
(**Adjective**).



Information Extraction

Finding relationships often involves locating a **Verb** between two **Named Entities**.



Parsing

Building syntax trees requires knowing if a word is the head of a phrase (**Verb**) or an argument (**Noun**).

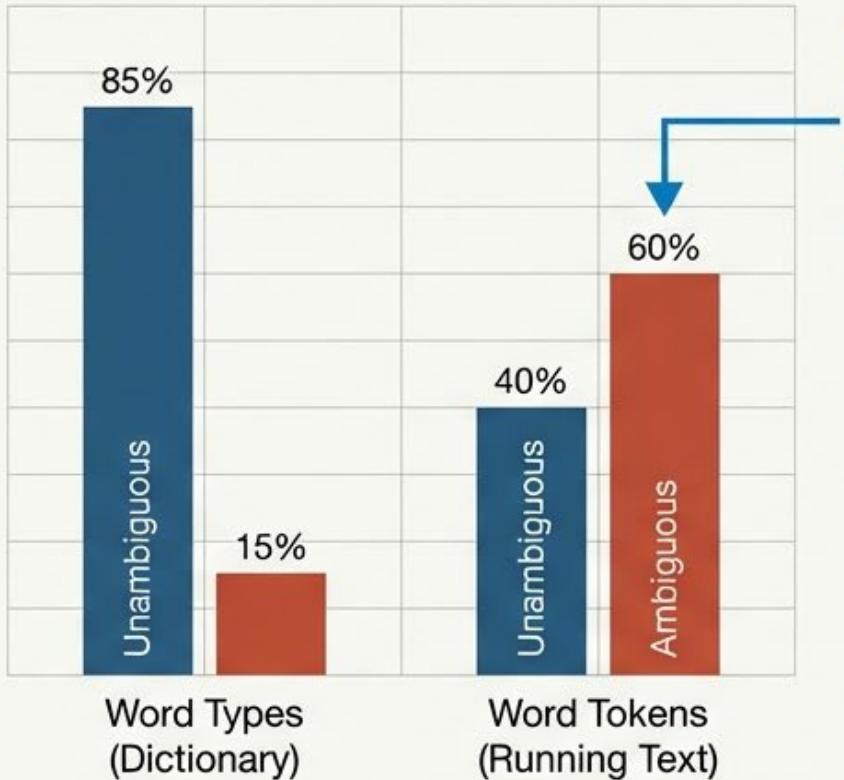


Sentiment Analysis

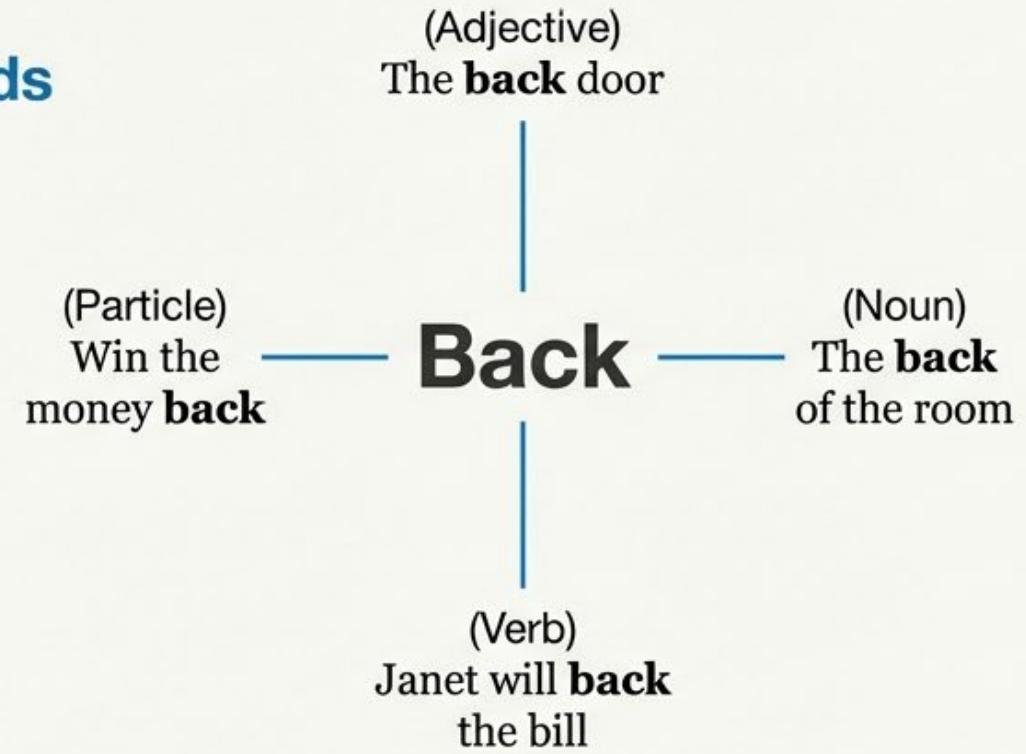
Adjectives (e.g., 'awesome', 'terrible') carry the bulk of sentiment weight.

How difficult is POS tagging in English?

Types vs. Tokens



Most
common words
are the most
ambiguous.



"Most words are simple, but the most common words are highly ambiguous."

Sources of information for POS tagging

Janet will back the bill
AUX/NOUN/VERB? NOUN/VERB?

Prior probabilities of word/tag

- "will" is usually an AUX

Identity of neighboring words

- "the" means the next word is probably not a verb

Morphology and wordshape:

- Prefixes unable: un- → ADJ
- Suffixes importantly: -ly → ADJ
- Capitalization Janet: CAP → PROPN

Named Entity Recognition (NER)

Named Entities

Part of speech tagging can tell us that words like **Janet, Stanford University, and Colorado** are all proper nouns; being a proper noun is a grammatical property of these words.

From a semantic perspective, these proper nouns refer to different kinds of **Named Entities**:

- **Janet** / Person
- **Stanford University** / Organization
- **Colorado** / Location

Often multi-word phrases

Type	Tag	Sample Categories	Example sentences
People	PER	people, characters	Turing is a giant of computer science.
Organization	ORG	companies, sports teams	The IPCC warned about the cyclone.
Location	LOC	regions, mountains, seas	Mt. Sanitas is in Sunshine Canyon .
Geo-Political Entity	GPE	countries, states	Palo Alto is raising the fees for parking.

Extended Named Entities

Named Entities are also extended to things that aren't entities:

- dates, times, money, events, products

Example

Citing high fuel prices, [ORG United Airlines] said [TIME Friday] it has increased fares by [MONEY \$6] per round trip on flights to some cities also served by lower-cost carriers. [ORG American Airlines], a unit of [ORG AMR Corp.], immediately matched the move, spokesman [PER Tim Wagner] said. [ORG United], a unit of [ORG UAL Corp.], said the increase took effect [TIME Thursday] and applies to most routes where it competes against discount carriers, such as [LOC Chicago] to [LOC Dallas] and [LOC Denver] to [LOC San Francisco].

What is Named Entity Recognition?

NER is the task of identifying and classifying named entities in text into predefined categories such as persons, organizations, locations, dates, and more.

At its core, NLP is just a two-step process:

- Span Detection (Identify Entity Spans): The system scans text to locate potential entities, such as proper nouns, dates, or numerical values, and determines where they start and end (e.g., recognizing that "New" and "York" form the entity "New York").
- Entity Classification (Tagging): The system assigns a category to the identified span based on pre-defined types, such as Person (PER), Organization (ORG), or Location (LOC).

Why NER matters?

NER is a prerequisite for many sophisticated NLP applications

The Bridge to Structured Information

"Microsoft acquired GitHub for \$7.5 billion. The deal was announced by CEO Satya Nadella in San Francisco."

{Organizations: Microsoft, GitHub ; **Person:** Satya Nadella ; **Location:** San Francisco ; **Money:** \$7.5 billion}

Question Answering

Input: Who founded Microsoft?

NER Helps: Extract PERSON entity associated with ORG 'Microsoft'

Relation Extraction

Input: Elon Musk is the CEO of Tesla"

NER Helps: Extract: (Elon Musk, CEO_OF, Tesla)

Document Summarization

Input: Long article about Apple's earnings

NER Helps: Focus on key entities: Apple, revenue, products

Why NER is hard

1. Segmentation

- In POS tagging, no segmentation problem since each word gets one tag.
- In NER, we have to find and segment the entities!

2. Type ambiguity

[PER Washington] was born into slavery on the farm of James Burroughs.

[ORG Washington] went up 2 games to 1 in the four-game series.

Blair arrived in [LOC Washington] for what may well be his last state visit.

In June, [GPE Washington] passed a primary seatbelt law.

BIO Tagging

How can we turn this structured problem into a sequence problem like POS tagging, with one label per word?

[PER Jane Villanueva] of [ORG United Airlines Holding] discussed the [LOC Chicago] route.

B: token that *begins* a span

I: tokens *inside* a span

O: tokens outside of any span

of tags (where n is #entity types):

- 1 O tag,
- n B tags,
- n I tags
- total of $2n+1$

Words	BIO Label
Jane	B-PER
Villanueva	I-PER
of	O
United	B-ORG
Airlines	I-ORG
Holding	I-ORG
discussed	O
the	O
Chicago	B-LOC
route	O
.	O

BIO Tagging variants: IO and BIOES

IO tagging, which loses some information by eliminating the B tag.

BIOES tagging, which adds an end tag E for the end of a span, and a span tag S for a span consisting of only one word.

Words	IO Label	BIO Label	BIOES Label
Jane	I-PER	B-PER	B-PER
Villanueva	I-PER	I-PER	E-PER
of	O	O	O
United	I-ORG	B-ORG	B-ORG
Airlines	I-ORG	I-ORG	I-ORG
Holding	I-ORG	I-ORG	E-ORG
discussed	O	O	O
the	O	O	O
Chicago	I-LOC	B-LOC	S-LOC
route	O	O	O
.	O	O	O

Standard algorithms for POS tagging and NER

Supervised Machine Learning given a human-labeled training set of text annotated with tags

- Hidden Markov Models
- Conditional Random Fields (CRF)
- Maximum Entropy Markov Models (MEMM)
- Neural sequence models (RNNs or Transformers)

All required a hand-labeled training set

Approach	Training Data	Accuracy	Complexity
Rule-Based	None	~85-90%	Low
HMM	Annotated corpus	~96-97%	Medium
CRF	Annotated corpus	~97%	Medium
RNNs	Large annotated corpus	~97-98%	High
Transformer	Massive pre-training	~98-99%	Very High

Hidden Markov Model for POS Tagging

Markov Chains: The Foundation

Markov Assumption

To predict the future, only the present matters. The past is irrelevant except through the current state.

Formula:

$$P(q_i = a \mid q_1 \dots q_{i-1}) = P(q_i = a \mid q_{i-1})$$

Think of it like predicting tomorrow's weather based ONLY on today's weather!

Markov Chain Components

Q - States

$$Q = \{q_1, q_2, \dots, q_N\}$$

The set of N possible states (e.g., HOT, COLD, WARM)

A - Transition Matrix

$$a_{ij} = P(q_j | q_i)$$

Probability of moving from state i to state j. Each row sums to 1.0

π - Initial Distribution

$$\pi = [\pi_1, \pi_2, \dots, \pi_N]$$

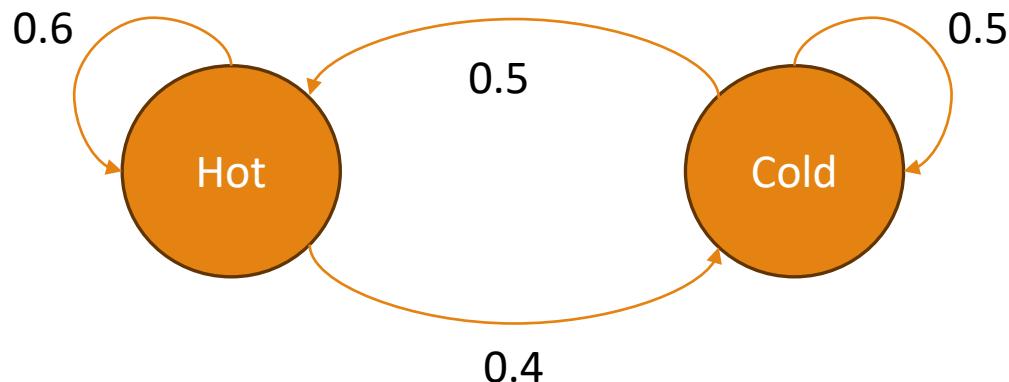
Probability of starting in each state. Sum equals 1.0

Weather Example

States: HOT, COLD

Transitions:

- HOT → HOT: 60%
- HOT → COLD: 40%
- COLD → COLD: 50%
- COLD → HOT: 50%



Transition Matrix

	HOT	COLD
HOT	0.6	0.4
COLD	0.5	0.5

If today is HOT, tomorrow has 60% chance of being HOT and 40% chance of being COLD.

Hidden Markov Model (HMM)

States are hidden, but observations are visible.

We only see OBSERVATIONS generated by hidden states.

Goal: infer the hidden states!

For example, we don't normally observe part-of-speech tags in a text.

Rather, we see words, and must infer the tags from the word sequence.

HMM Components

B - Emission Probabilities

Probability of observing o_t given state q_i

$$P(3 \text{ ice creams} | \text{HOT}) = 0.4$$

$$b_i(o_t) = P(o_t | q_i)$$

O - Observation Sequence

The sequence we actually see

$$O = \{3, 1, 3\} \text{ (ice cream counts)}$$

$$O = \{o_1, o_2, \dots, o_T\}$$

Plus all Markov Chain components: Q (states), A (transitions), π (initial distribution)

Weather HMM Example

Transition Probabilities (Hidden States)

	HOT	COLD
HOT	0.7	0.3
COLD	0.4	0.6

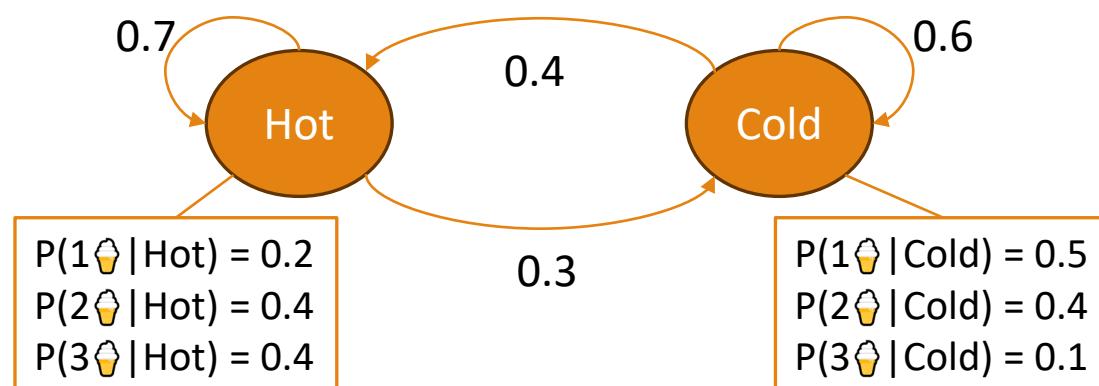
Emission Probabilities (Observations)

	1 	2 	3 
HOT	0.2	0.4	0.4
COLD	0.5	0.4	0.1

Initial Probabilities: $P(HOT) = 0.8$ $P(COLD) = 0.2$

Interpretation

- If it's HOT today, 70% chance it stays HOT tomorrow
 - If it's HOT, there's 40% chance we observe 3 ice creams sold
 - We start with 80% probability that day 1 is HOT



HMM: Two Simplifying Assumptions

1

Markov Assumption

$$P(q_i \mid q_1 \dots q_{i-1}) = P(q_i \mid q_{i-1})$$

Current state depends only on previous state, not entire history

2

Output Independence

$$P(o_i \mid q_1 \dots q_T, o_1 \dots o_T) = P(o_i \mid q_i)$$

Observation depends only on current state, not others

The components of an HMM tagger

An HMM tagger has two components, the **A (Transition)** and **B (Emission)** probabilities.

Transition Matrix

The **A matrix** contains the tag transition probabilities:

$$P(t_i|t_{i-1}) = \frac{C(t_{i-1}, t_i)}{C(t_{i-1})}$$

$C(t_{i-1}, t_i)$: how often (count) t_{i-1} is followed by t_i in the labeled corpus.

$C(t_{i-1})$: how often (count) t_{i-1} found in the labeled corpus.

Emission Matrix

The **B matrix** represent the probability, given a tag, that it will be associated with a given word

$$P(w_i|t_i) = \frac{C(t_i, w_i)}{C(t_i)}$$

$C(t_i, w_i)$: how often (count) t_i is associated with w_i in the labeled corpus.

$C(t_i)$: how often (count) t_i found in the labeled corpus.

HMM tagging as decoding

Decoding

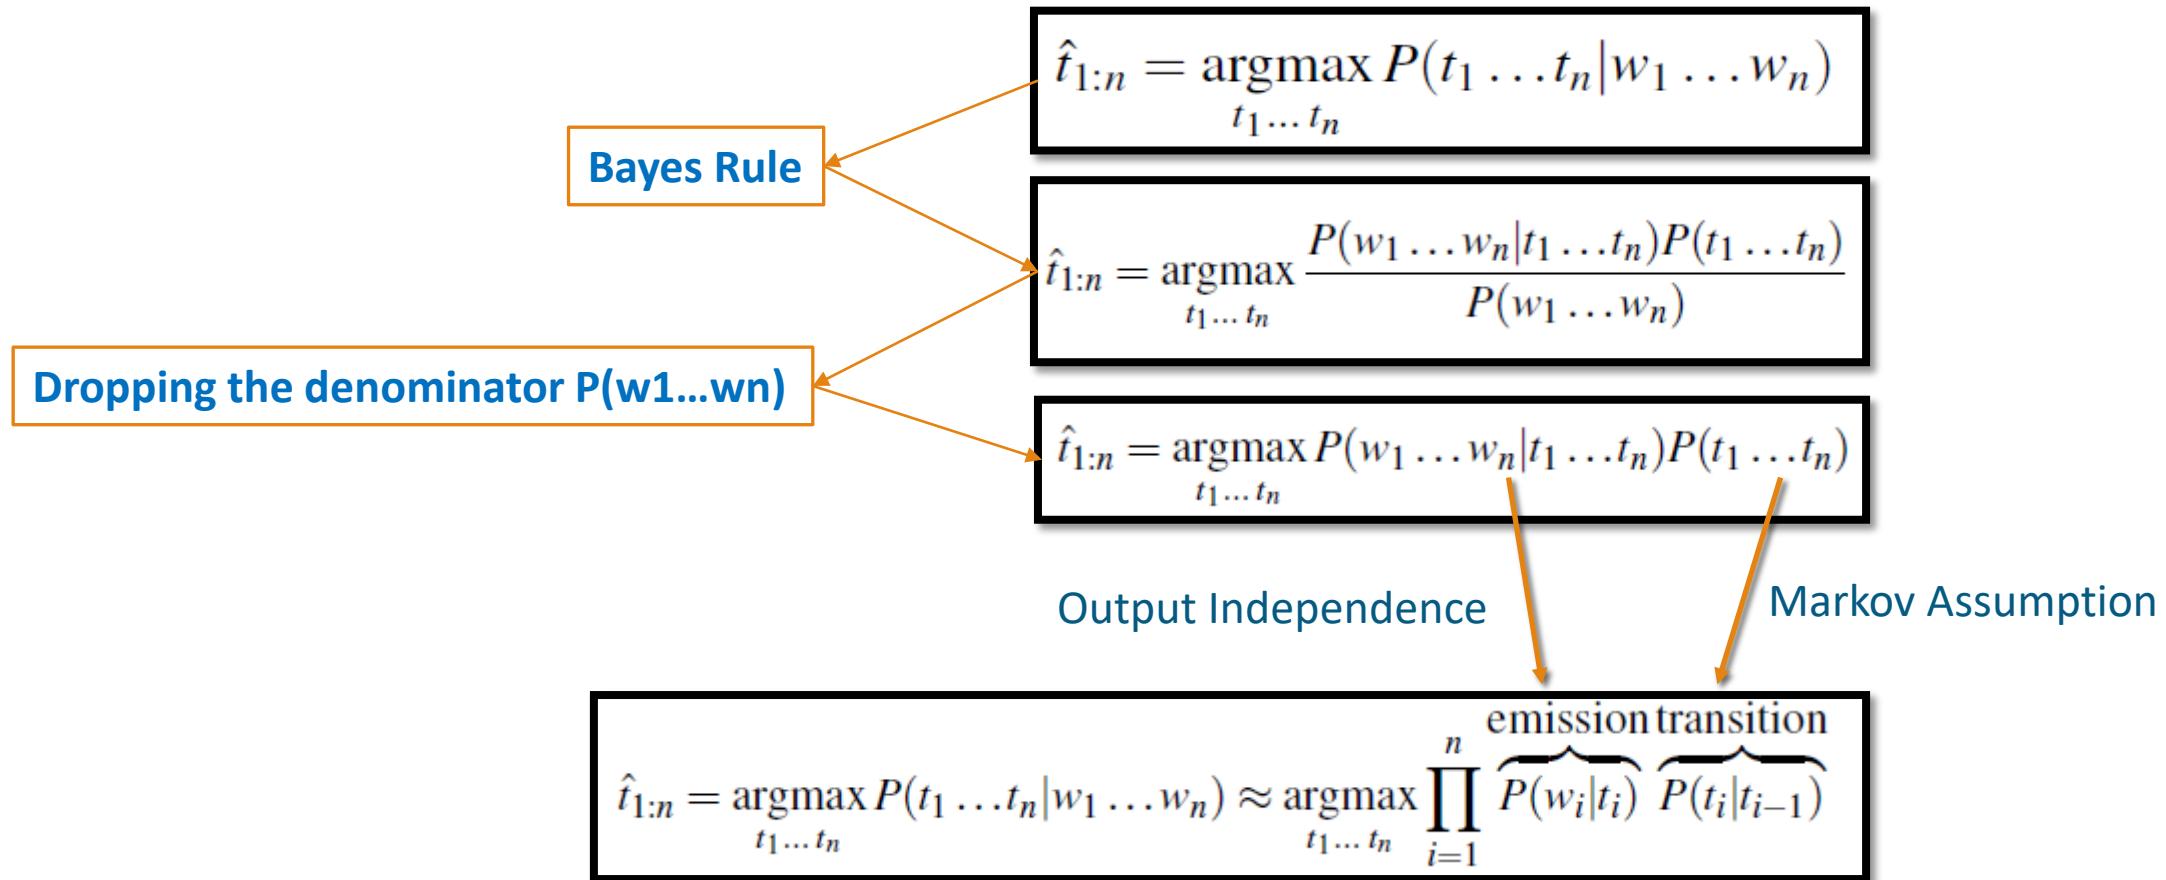
Given as input an HMM $\lambda = (A, B)$ and a sequence of observations $O = o_1, o_2, \dots, o_T$, find the most probable sequence of states $Q = q_1 q_2 q_3 \dots q_T$

HMM Decoding for POS Tagging

$$\hat{t}_{1:n} = \underset{t_1 \dots t_n}{\operatorname{argmax}} P(t_1 \dots t_n | w_1 \dots w_n)$$

choose the tag sequence $t_1 \dots t_n$ that is most probable given the observation sequence of n words $w_1 \dots w_n$

Simplifying Decoding



Decoding Algorithms

Brute Force, $O(N^T)$

N states (POS tags) and T observations (word sequence length)

Viterbi Algorithm, $O(T * N^2)$

Dynamic programming algorithm to find the most likely sequence of hidden states

Viterbi Algorithm

Key Idea

Build up the best path incrementally by storing the maximum probability of reaching each state at each time step, along with the path that led there.

Algo Steps

1. Initialization $v1(s) = \pi_s \times bs(o_1)$, for all s

Start probability \times emission for first observation

2. Recursion $vt(j) = \text{MAX}_i vt-1(i) \times a_{ij} \times bj(ot)$

Take MAX, store backpointer

3. Termination $P^* = \text{MAX}_j vT(j)$

Highest final probability

4. Backtrace *Follow backpointers*

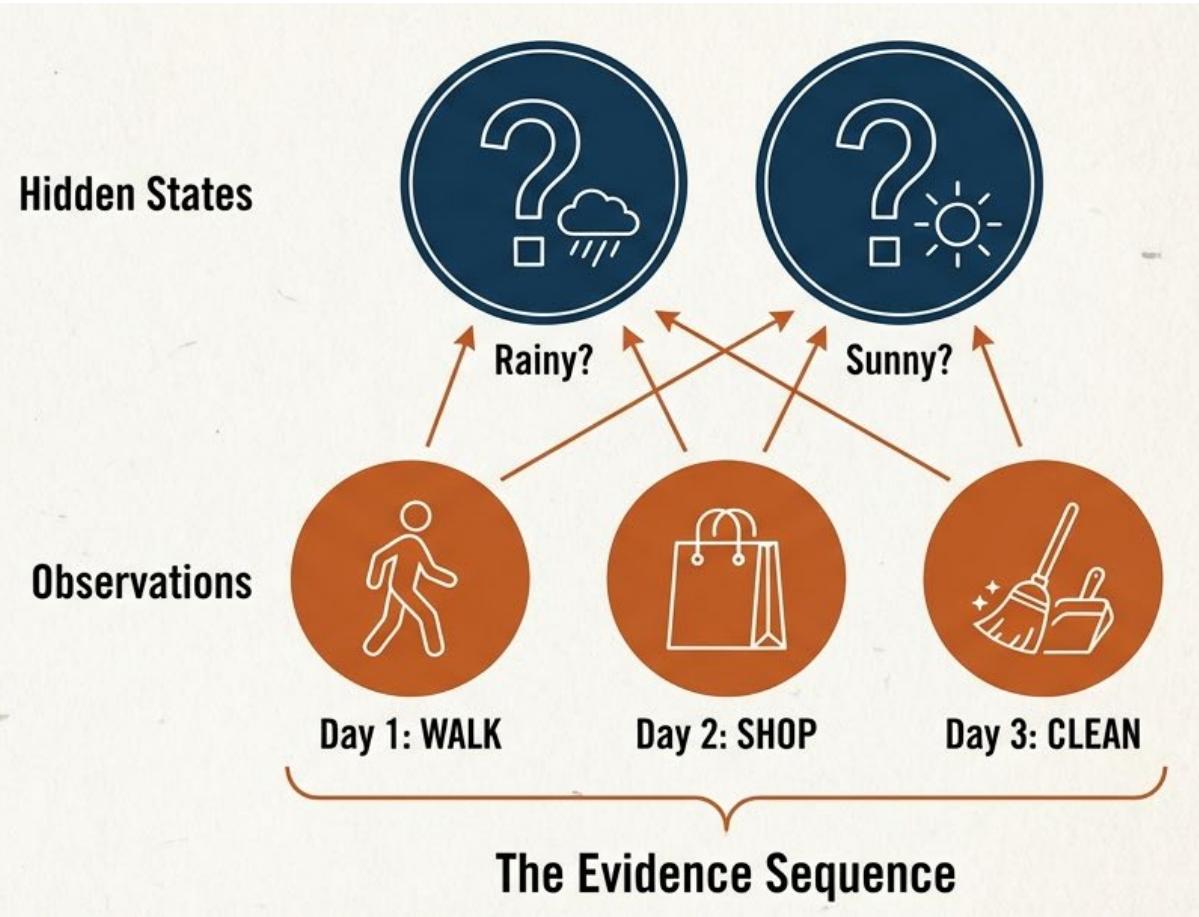
Reconstruct best sequence

Viterbi Algorithm

```
function VITERBI(observations of len  $T$ ,state-graph of len  $N$ ) returns best-path, path-prob
    create a path probability matrix  $viterbi[N,T]$ 
    create a backpointer matrix  $backpointer[N,T]$ 
    for each state  $s$  from 1 to  $N$  do ; initialization step
         $viterbi[s,1] \leftarrow \pi_s * b_s(o_1)$ 
         $backpointer[s,1] \leftarrow 0$ 
    for each time step  $t$  from 2 to  $T$  do ; recursion step
        for each state  $s$  from 1 to  $N$  do
             $viterbi[s,t] \leftarrow \max_{s'=1}^N viterbi[s',t-1] * a_{s',s} * b_s(o_t)$ 
             $backpointer[s,t] \leftarrow \operatorname{argmax}_{s'=1}^N viterbi[s',t-1] * a_{s',s} * b_s(o_t)$ 
     $bestpathprob \leftarrow \max_{s=1}^N viterbi[s,T]$  ; termination step
     $bestpathpointer \leftarrow \operatorname{argmax}_{s=1}^N viterbi[s,T]$  ; termination step
     $bestpath \leftarrow$  the path starting at state  $bestpathpointer$ , that follows  $backpointer[]$  to states back in time
    return bestpath, bestpathprob
```

Viterbi Algorithm in Action: Scenario

The Scenario:
We have a friend, Bob.
We never see the
weather where he lives
(Hidden States), but
we see his daily blog
posts about his
activities
(Observations).



Viterbi Algorithm in Action: Given

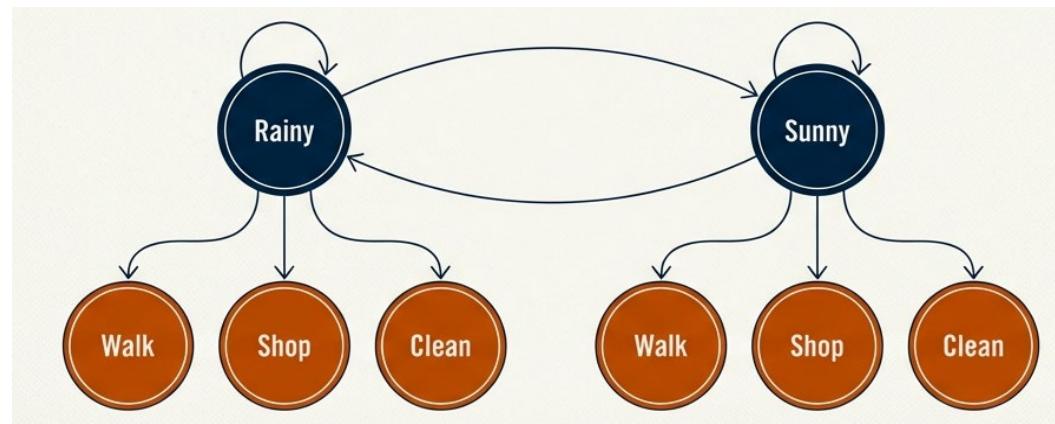
Transition Probabilities (Hidden States)

	Rainy	Sunny
Rainy	0.7	0.3
Sunny	0.4	0.6

Emission Probabilities (Observations)

	Walk	Shop	Clean
Rainy	0.1	0.4	0.5
Sunny	0.6	0.3	0.1

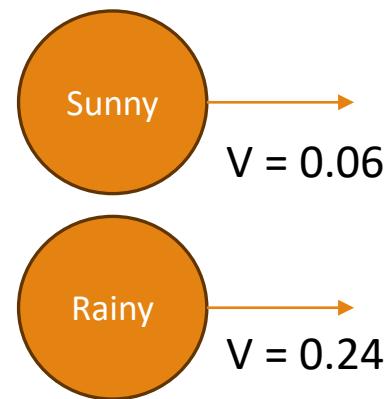
Initial Probabilities: $P(\text{Rainy}) = 0.6$ $P(\text{Sunny}) = 0.4$



Viterbi Algorithm in Action: Initialization

DAY 1: Observation Walk

- $v1(\text{Rainy}) = P(\text{Rainy}) * P(\text{Rainy} | \text{Walk}) = 0.6 * 0.1 = 0.06$
- $V1(\text{Sunny}) = P(\text{Sunny}) * P(\text{Sunny} | \text{Walk}) = 0.4 * 0.6 = 0.24$



Day 1: Walk

Viterbi Algorithm in Action: Recursion

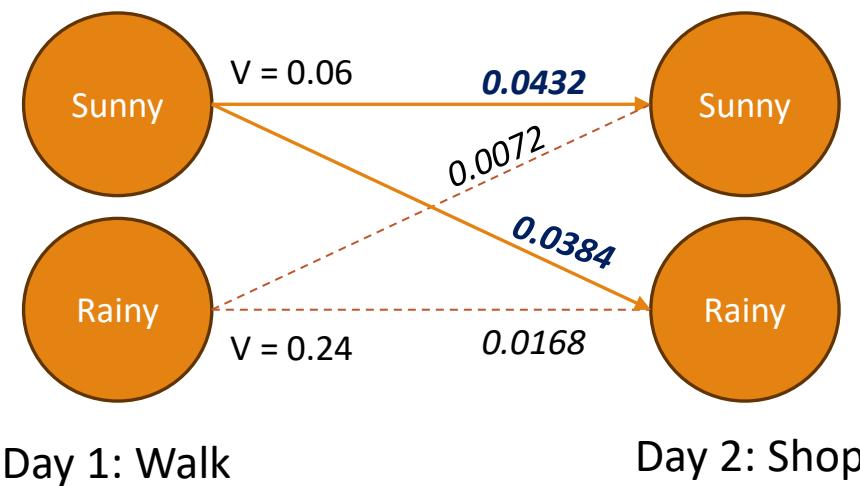
DAY 2: Observation Shop

□ $v2(\text{Rainy}) = \text{Max} (v1(\text{Rainy}) * P(\text{Rainy}/\text{Shop}) * P(\text{Rainy}/\text{Rainy}), v1(\text{Sunny}) * P(\text{Rainy}/\text{Shop}) * P(\text{Rainy}/\text{Sunny}))$

$$v2(\text{Rainy}) = \text{Max} (0.06 * 0.4 * 0.7 = 0.0168, 0.24 * 0.4 * 0.4 = \mathbf{0.0384}) \quad \boxed{\text{Backpointer}[\text{Rainy}, \text{Day 1}] = \text{Sunny}}$$

□ $v2(\text{Sunny}) = \text{Max} (v1(\text{Rainy}) * P(\text{Sunny}/\text{Shop}) * P(\text{Sunny}/\text{Rainy}), v1(\text{Sunny}) * P(\text{Sunny}/\text{Shop}) * P(\text{Sunny}/\text{Sunny}))$

$$v2(\text{Sunny}) = \text{Max} (0.06 * 0.3 * 0.4 = 0.0072, 0.24 * 0.3 * 0.6 = \mathbf{0.0432}) \quad \boxed{\text{Backpointer}[\text{Sunny}, \text{Day 1}] = \text{Sunny}}$$



Viterbi Algorithm in Action: Recursion

DAY 3: Observation Clean

□ $v_3(\text{Rainy}) = \text{Max} (v_2(\text{Rainy}) * P(\text{Rainy}/\text{Clean}) * P(\text{Rainy}/\text{Rainy}), v_2(\text{Sunny}) * P(\text{Rainy}/\text{Clean}) * P(\text{Rainy}/\text{Sunny}))$

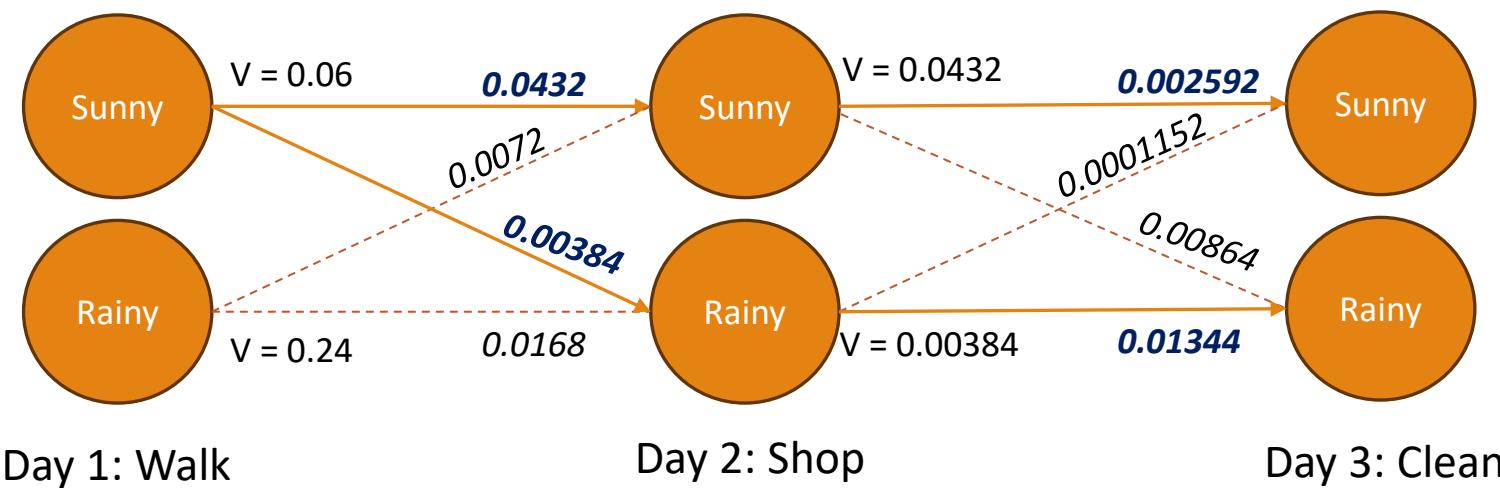
$$v_3(\text{Rainy}) = \text{Max} (0.0384 * 0.5 * 0.7 = \mathbf{0.01344}, 0.0432 * 0.5 * 0.4 = 0.00864)$$

Backpointer[Rainy, Day 2] = Rainy

□ $v_3(\text{Sunny}) = \text{Max} (v_2(\text{Rainy}) * P(\text{Sunny}/\text{Clean}) * P(\text{Sunny}/\text{Rainy}), v_2(\text{Sunny}) * P(\text{Sunny}/\text{Clean}) * P(\text{Sunny}/\text{Sunny}))$

$$v_3(\text{Sunny}) = \text{Max} (0.00384 * 0.1 * 0.3 = 0.0001152, 0.0432 * 0.1 * 0.6 = \mathbf{0.002592})$$

Backpointer[Sunny, Day 2] = Sunny



Viterbi Algorithm in Action: Termination and Backtrace

Termination

$P = \text{Max}(v_3(\text{Sunny}), v_3(\text{Rainy})) = \text{Max}(0.002592, \mathbf{0.01344}) = \mathbf{0.01344} \rightarrow \text{Last Hidden State is Rainy}$

Backtrace

