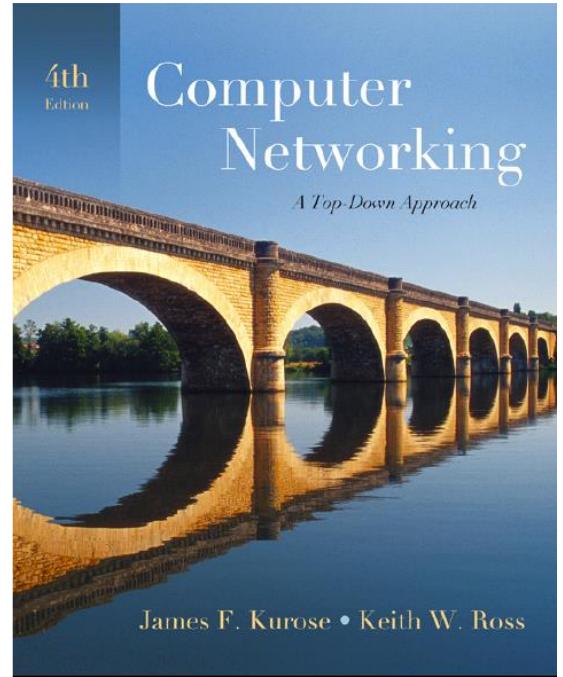


# Chapter 2

## Application Layer



*Computer Networking:  
A Top Down Approach,  
Jim Kurose, Keith Ross  
Addison-Wesley.*

# Chapter 2: Application layer

- 2.1 Principles of network applications
- 2.2 Web and HTTP
- 2.3 FTP
- 2.4 Electronic Mail
  - ❖ SMTP, POP3, IMAP
- 2.5 DNS
- 2.6 P2P Applications
- 2.7 Socket programming with TCP
- 2.8 Socket programming with UDP

# Chapter 2: Application Layer

## Our goals:

- conceptual, architectural aspects of network application protocols
  - ❖ transport-layer service models
  - ❖ client-server paradigm
  - ❖ peer-to-peer paradigm
- learn about protocols
  - ❖ HTTP
  - ❖ FTP
  - ❖ SMTP / POP3 / IMAP
  - ❖ DNS

# Some network apps

- e-mail
- web
- instant messaging
- remote login
- P2P file sharing
- multi-user network games
- streaming stored video clips
- voice over IP
- real-time video conferencing
- grid computing
- ...
- 
- 

Note: different applications may have different

- Requirements (delay, loss, Tput, jitter bounds, security)
- Number of participants (unicast, multicast, broadcast, manycast, profilecast)
- Architecture (client-server, p2p, flat, hierarchical, hybrid, self-configuring)

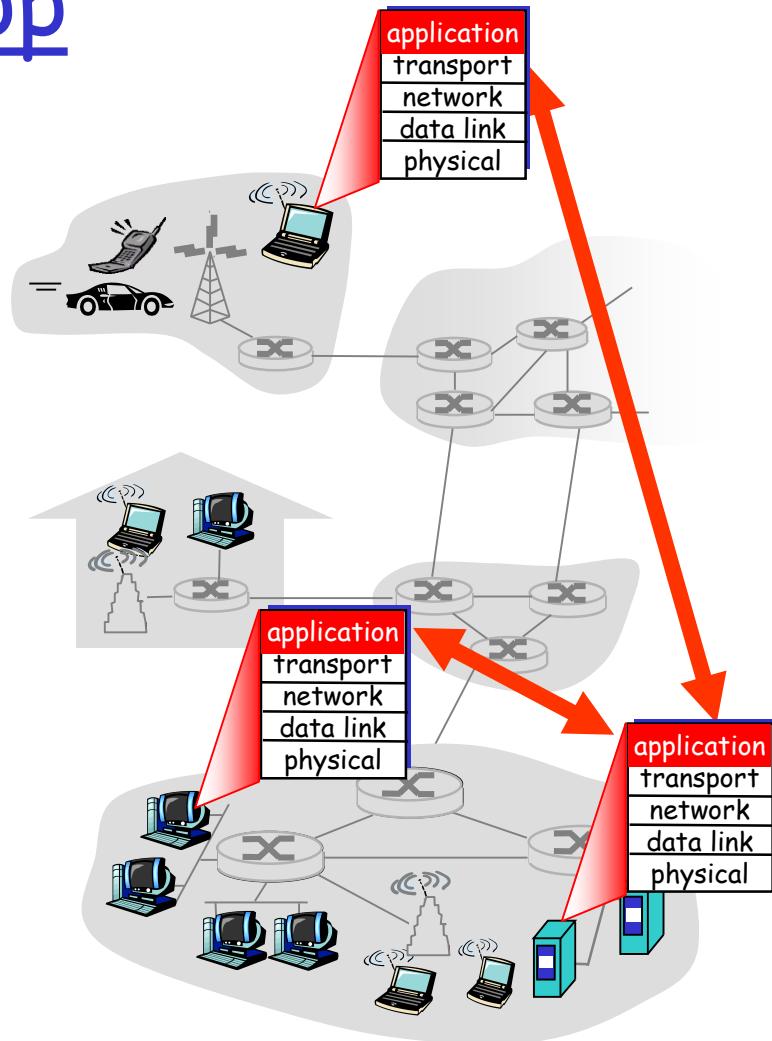
# Creating a network app

## Application programs

- ❖ run on *end systems*
- ❖ communicate over network

little software written for devices in network core

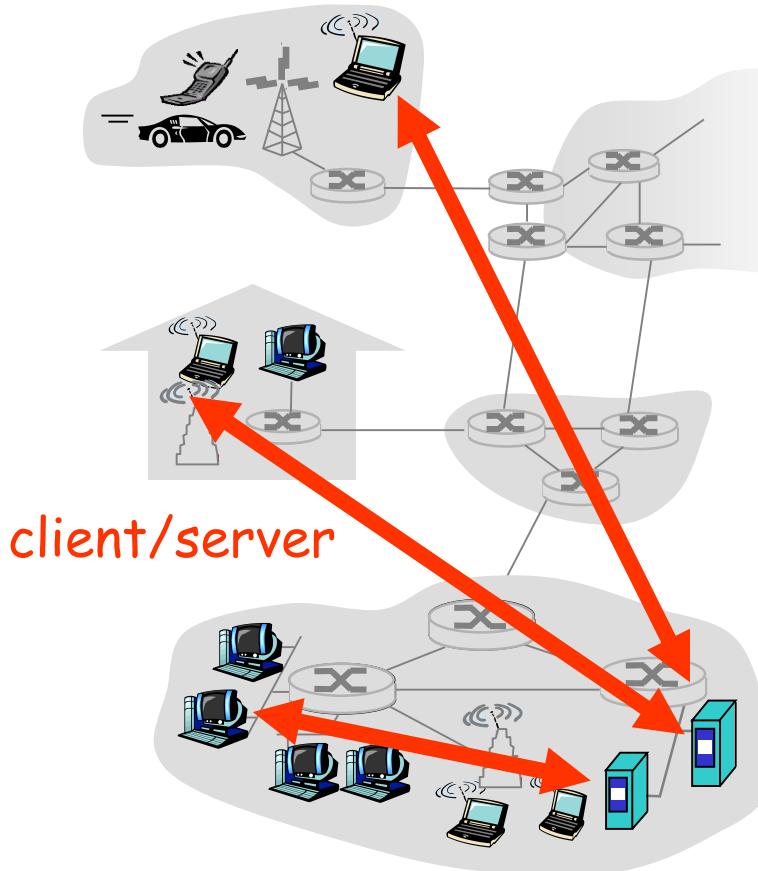
- ❖ network core devices do not run user applications
- ❖ applications on end systems allows for rapid app development, propagation



# Application architectures

- Client-server
- Peer-to-peer (P2P)
- Hybrid of client-server and P2P

# Client-server architecture



**server:**

- ❖ always-on host
- ❖ permanent IP address
- ❖ server farms for scaling

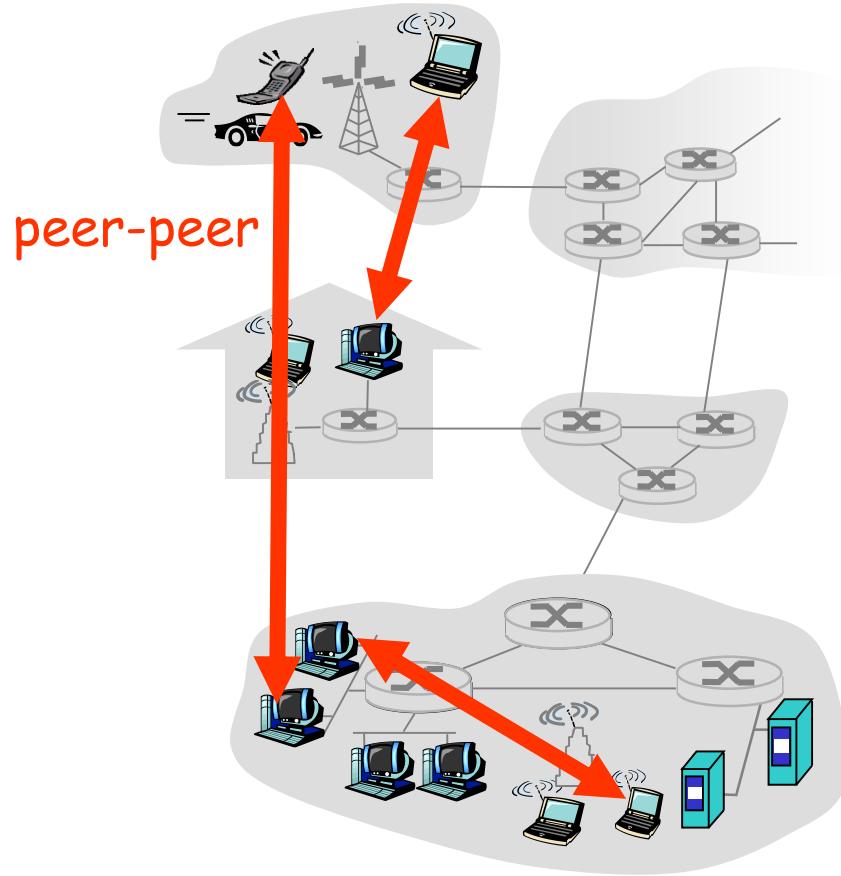
**Clients (in general):**

- ❖ communicate with server
- ❖ intermittently connected
- ❖ have dynamic IP addresses
- ❖ do not communicate directly with each other

# Pure P2P architecture

- No 'always-on' server
- arbitrary end systems directly communicate
- peers intermittently connected & change IP addresses
- example: Gnutella

Highly scalable but difficult to manage



# Hybrid of client-server and P2P

## Skype

- ❖ voice-over-IP P2P application
- ❖ centralized server: finding address of remote party
- ❖ client-client connection: direct (not through server)

## Instant messaging

- ❖ chatting between two users is P2P
- ❖ centralized service: client detection & location
  - user registers IP address with central server
  - uses central server to find addresses of buddies

# Processes communicating

**Process:** program running within a host.

- within same host, two processes communicate using **inter-process communication** (defined by OS).
- processes in different hosts communicate by exchanging **messages**

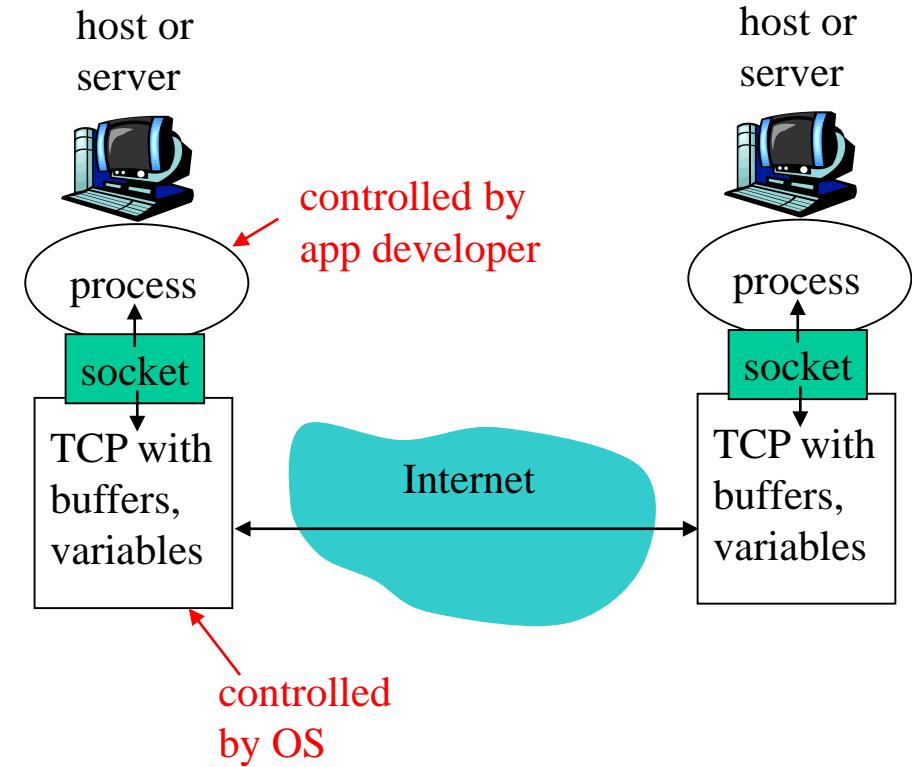
**Client process:** process that initiates communication

**Server process:** process that waits to be contacted

- Note: applications with P2P architectures have client processes & server processes

# Sockets

- ❑ process sends/receives messages to/from its **socket**



- ❑ API: (1) choice of transport protocol; (2) ability to fix a few parameters (**more on this later**)

# Addressing processes

- **Q:** does IP address of host on which process runs suffice for identifying the process?
  - ❖ **A:** No, many processes can run on same host
- **identifier** includes **IP address & port numbers** associated with process on host.
- Example port numbers:
  - ❖ HTTP server: 80
  - ❖ Mail server: 25
- to send HTTP message to gaia.cs.umass.edu web server:
  - ❖ **IP address:** 128.119.245.12
  - ❖ **Port number:** 80
- more later...

# What transport service does an app need?

## Data loss

- some apps (e.g., audio) can tolerate some loss (~10%)
- others (e.g., FTP, telnet) require 100% delivery

## Timing

- some apps (e.g., VoIP, interactive games) require low (bounded) delay (and/or jitter) to be "effective"
  - ❖ Example: VoIP jitter or handoff delay bound is ~200ms
- others (e.g., FTP) are tolerant to some delay/jitter
  - ❖ Some multimedia apps use buffering & playback point adjustment

## Bandwidth

- some apps (e.g., multimedia) require minimum amount of bandwidth to be "effective"
- other apps ("elastic apps") make use of whatever bandwidth they get

# Transport service requirements of common apps

<b>Application</b>	<b>Data loss</b>	<b>Bandwidth</b>	<b>Time Sensitive</b>
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5kbps-1Mbps video:10kbps-5Mbps	yes, 100's msec
stored audio/video	loss-tolerant	same as above	yes, few secs
interactive games	loss-tolerant	few kbps up	yes, 100's msec
instant messaging	no loss	elastic	yes and no

# Internet transport protocols services

## TCP service:

- connection-oriented:** setup required between client and server processes
- reliable transport** between sending and receiving process
- flow control:** sender won't overwhelm receiver
- congestion control:** throttle sender when network overloaded
- does not provide:** timing, minimum bandwidth guarantees, multicast support

## UDP service:

-

# Internet apps: application, transport protocols

Application	Application layer protocol	Underlying transport protocol
e-mail	SMTP [RFC 2821]	TCP
remote terminal access	Telnet [RFC 854]	TCP
Web	HTTP [RFC 2616]	TCP
file transfer	FTP [RFC 959]	TCP
streaming multimedia	proprietary (e.g. RealNetworks)	TCP or UDP
Internet telephony	proprietary (e.g., Vonage,Dialpad)	typically UDP

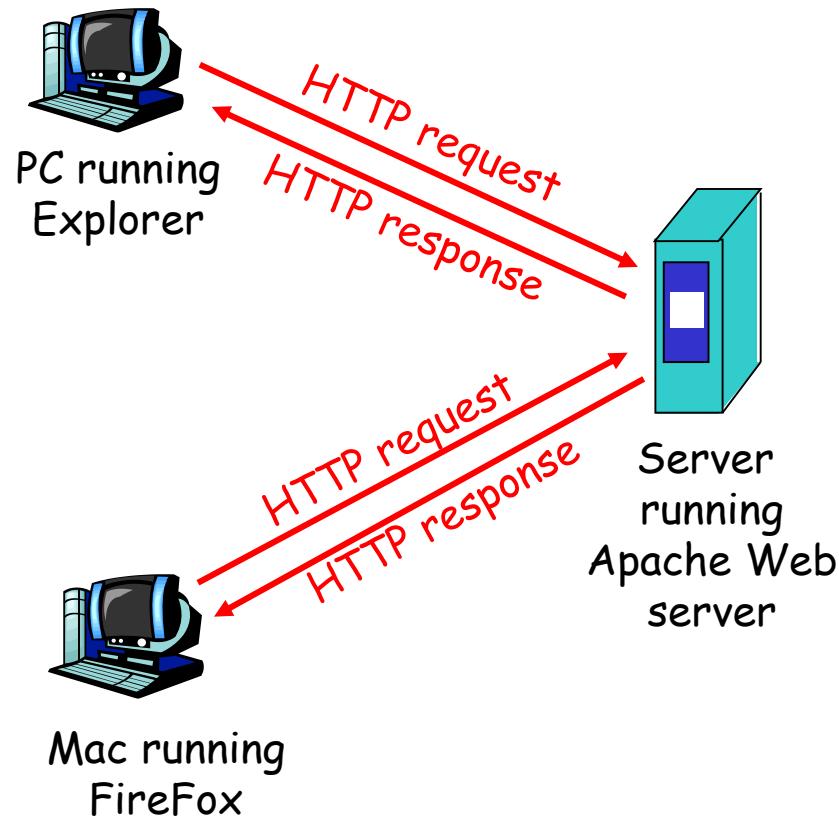
# Chapter 2: Application layer

- 2.1 Principles of network applications
  - ❖ app architectures
  - ❖ app requirements
- 2.2 Web and HTTP
- 2.4 Electronic Mail
  - ❖ SMTP, POP3, IMAP
- 2.5 DNS
- 2.6 P2P file sharing
- 2.7 Socket programming with TCP
- 2.8 Socket programming with UDP

# HTTP overview

## HTTP: hypertext transfer protocol

- client/server model
  - ❖ *client*: browser to request & receive Web objects
  - ❖ *server*: Web server sends objects in response to requests
- HTTP 1.0: RFC 1945
- HTTP 1.1: RFC 2068 (persistent TCP)



# HTTP overview (continued)

## Uses TCP:

- 1. client initiates TCP connection to server, port 80
- 2. server accepts TCP connection from client
- 3. HTTP (application-layer) messages exchanged between HTTP client and HTTP server
- 4. TCP connection closed

A 'state' is information kept in memory of a host, server or router to reflect past events: such as routing tables, data structures or database entries

## HTTP is "stateless"

- server maintains no information about past client requests

Protocols that maintain "state" are complex!

- history (state) is maintained
- if server/client crashes, views of "state" may be inconsistent, must be reconciled
- state is added via 'cookies'

## Design Issues:

- Stateful vs Stateless
- Hard State vs Soft State

# HTTP connections

## I. Nonpersistent HTTP

- At most one object is sent over a TCP connection. Used in HTTP/1.0

## II. Persistent HTTP

- Multiple objects can be sent over single TCP connection. Used in HTTP/1.1 by default:
  - ❖ A. persistent with pipelining
  - ❖ B. persistent without pipelining

# I. Nonpersistent HTTP

User enters URL

someSchool.edu/someDepartment/home.index

(contains text,  
references to 10  
jpeg images)

1a. HTTP client initiates TCP connection  
to someSchool.edu : port 80

1b. HTTP server @ someSchool.edu  
port 80. "accepts" connection,  
notifying client

2. HTTP client sends *request message* (containing URL) into TCP  
connection. Message indicates  
that client wants object  
someDepartment/home.index

3. HTTP server receives request  
message, sends *response message*  
containing requested object

time  
↓

# I. Nonpersistent HTTP (cont.)

time  
↓

4. HTTP server closes TCP connection.
5. HTTP client receives response message containing html file, displays html. Parsing html file, finds 10 referenced jpeg objects
6. Steps 1-5 repeated for each of 10 jpeg objects

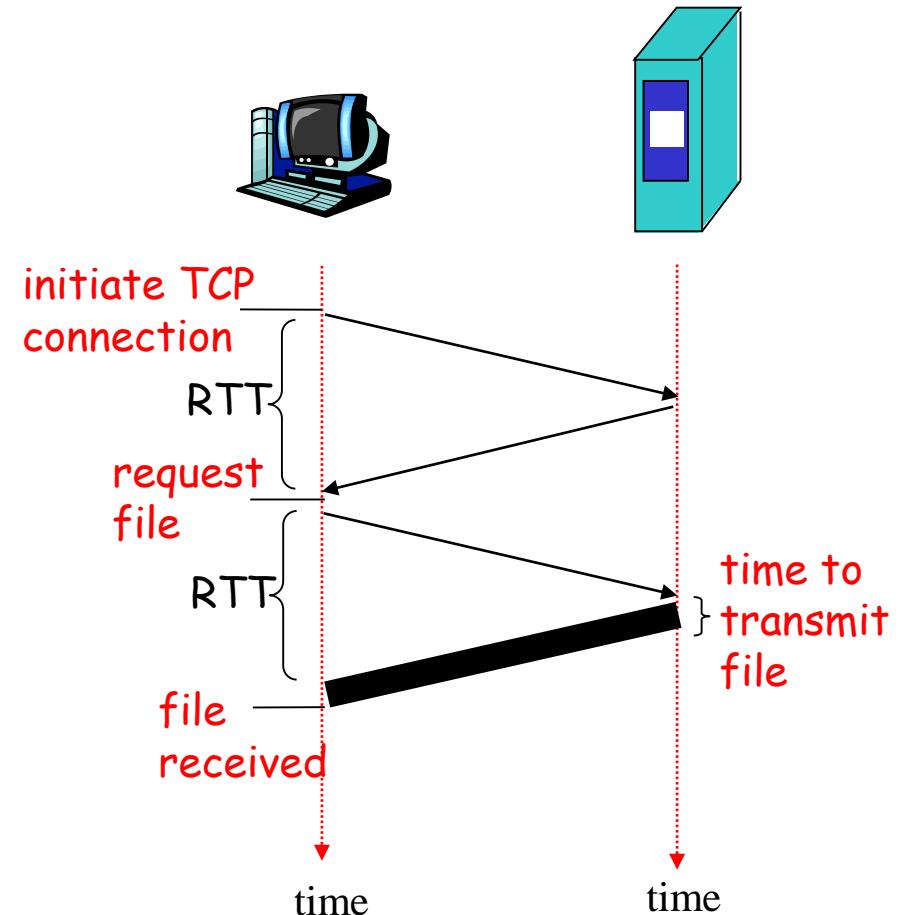
# I. Non-Persistent HTTP: Response time

**Definition of RTT:** time to send request from client to server and back.

**Response time:**

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- file transmission time

$$\text{total} = 2\text{RTT} + \text{transmit time}$$



# Persistent HTTP

## I. Nonpersistent HTTP issues:

- requires 2 RTTs per object
- OS overhead for each TCP connection
- browsers often open parallel TCP connections to fetch referenced objects

## II. Persistent HTTP

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection

### A. Persistent without pipelining:

- client issues new request only when previous response has been received
- one RTT for each referenced object

### B. Persistent with pipelining:

- default in HTTP/1.1
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects

# Web caches (proxy servers)

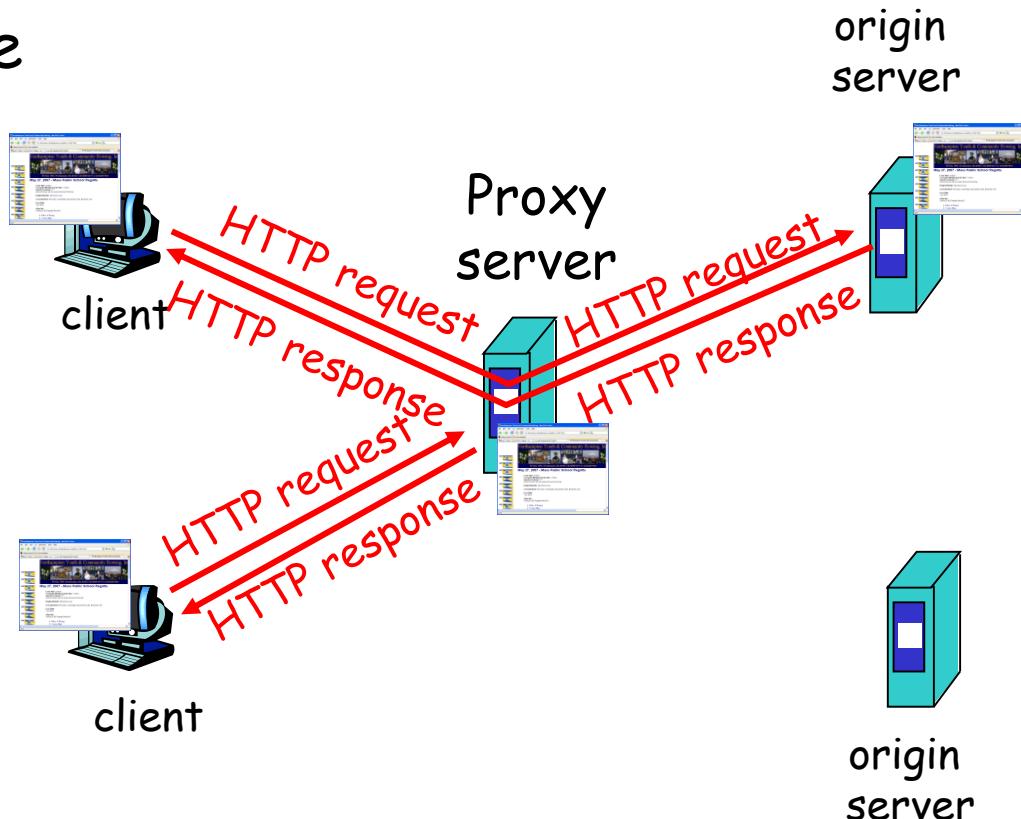
**Goal:** satisfy client request without involving origin server

- user sets browser:

Web accesses via cache

- browser sends all HTTP requests to cache

- ❖ object in cache: cache returns object
  - ❖ else cache requests object from origin server, then returns object to client
  - ❖ Cache keeps copy of object for future use



- Can all objects be cached?
- Proxy vs. local browser cache

# More about Web caching

- cache acts as both client and server
- typically cache is installed by ISP (university, company, residential ISP)

## Why Web caching?

- 1. reduce response time for client request
- 2. reduce traffic on an institution's access link.

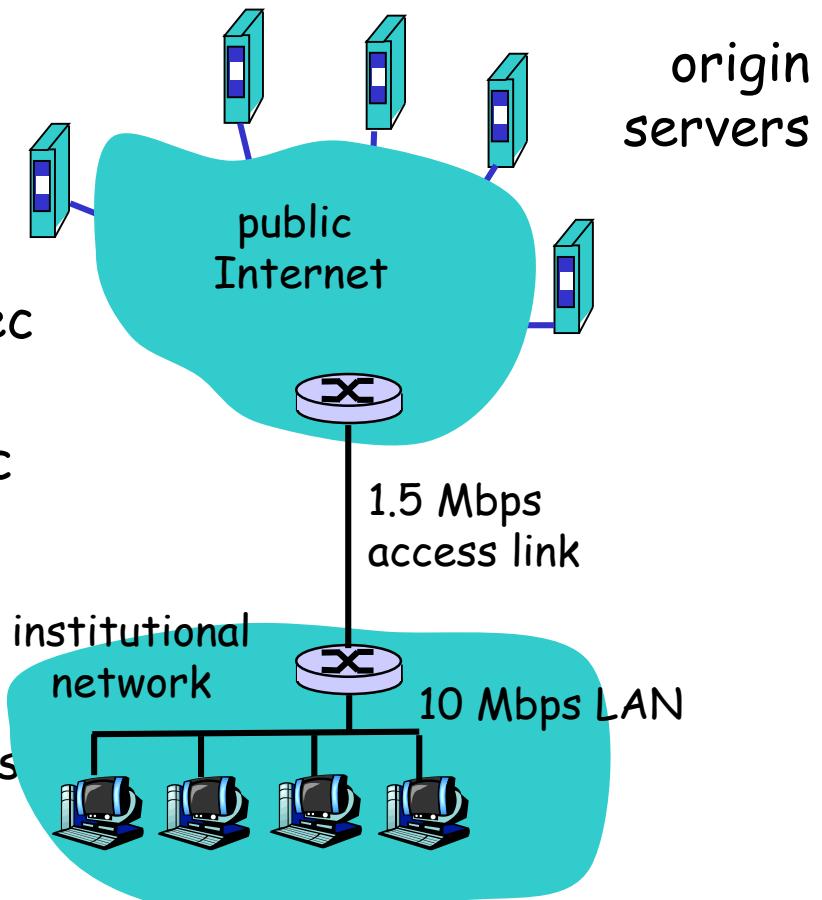
# Caching example

## Assumptions

- average object size = 100k bits
- avg. request rate from institution's browsers = 15 req/sec
- delay from institutional router to any origin server and back = 2 sec

## Consequences

- utilization on LAN = 15%
- utilization on access link = 100%
- total delay = Internet delay + access delay + LAN delay  
= 2 sec + minutes + milliseconds



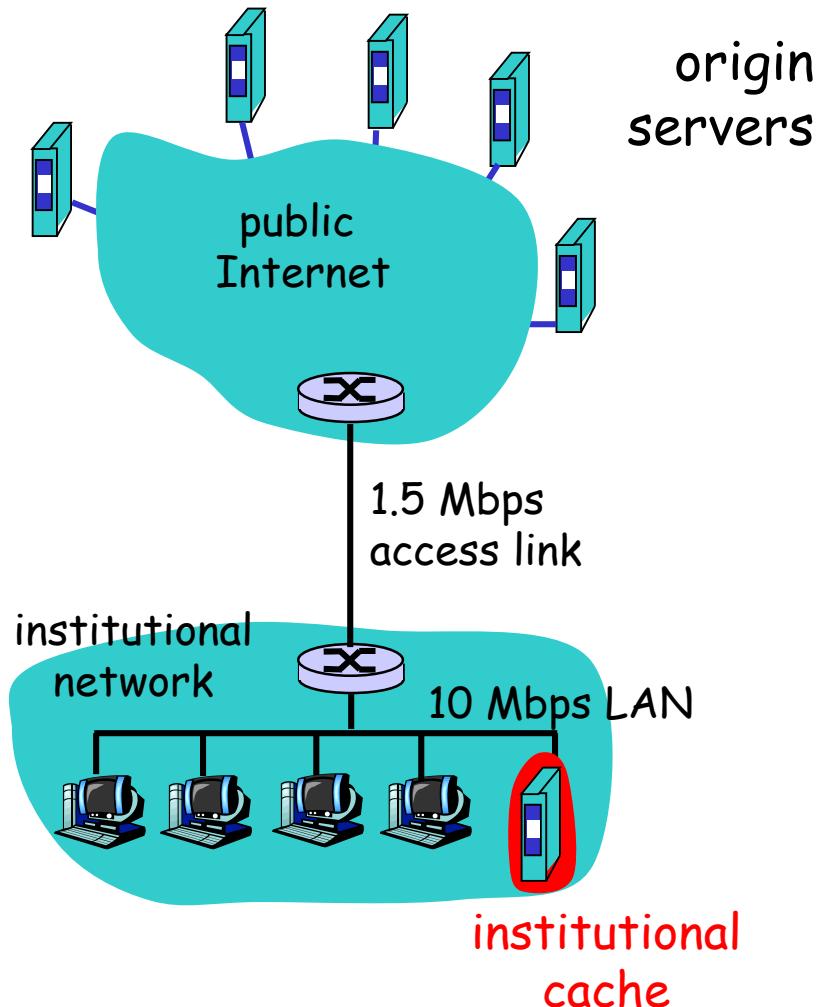
# Caching example (cont)

## one solution: install cache

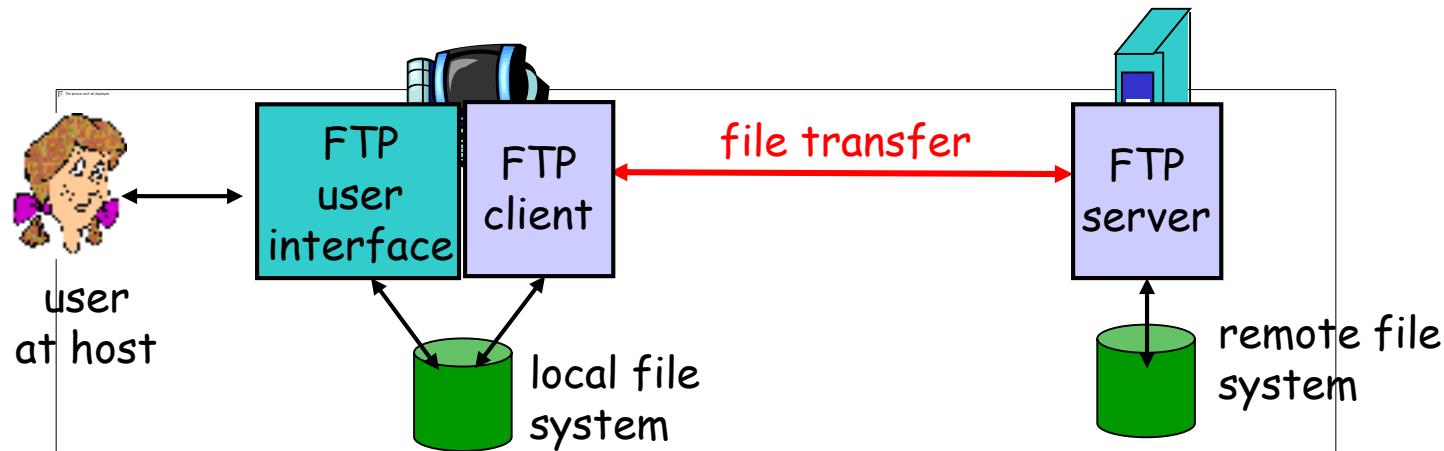
- suppose hit rate is 0.4

## consequence

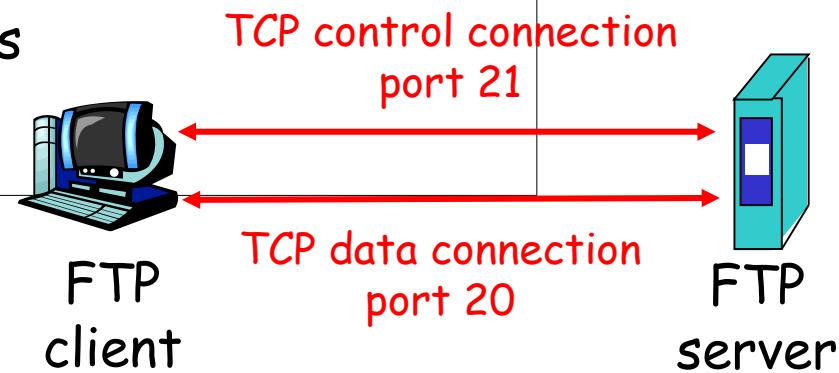
- 40% requests will be satisfied almost immediately
- 60% requests satisfied by origin server
- utilization of access link reduced to 60%, resulting in negligible delays (say 10 msec)
- total avg delay = Internet delay + access delay + LAN delay =  $.6 * (2.01)$  secs +  $.4 * \text{milliseconds} < 1.4$  secs



# FTP: the file transfer protocol

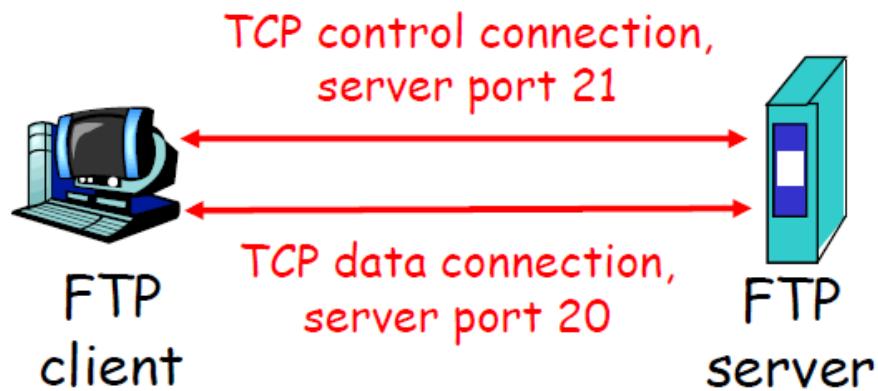


- client/server model
  - ❖ *client*: side initiating transfer, *server*: remote host
- ftp: RFC 959, ftp server: port 21
- Separate data and control connections
- FTP server maintains "state":
  - ❖ current directory, earlier authentication



# FTP: separate control, data connections

- ❖ FTP client contacts FTP server at port 21, TCP is transport protocol
- ❖ client authorized over control connection
- ❖ client browses remote directory by sending commands over control connection.
- ❖ when server receives file transfer command, server opens 2<sup>nd</sup> TCP connection (for file) to client
- ❖ after transferring one file, server closes data connection.



- ❖ server opens another TCP data connection to transfer another file.
- ❖ control connection: "**out of band**"
- ❖ FTP server maintains "state": current directory, earlier authentication

# FTP commands, responses

## sample commands:

- ❖ sent as ASCII text over control channel
- ❖ **USER *username***
- ❖ **PASS *password***
- ❖ **LIST** return list of file in current directory
- ❖ **RETR *filename*** retrieves (gets) file
- ❖ **STOR *filename*** stores (puts) file onto remote host

## sample return codes

- ❖ status code and phrase (as in HTTP)
- ❖ **331 Username OK, password required**
- ❖ **125 data connection already open; transfer starting**
- ❖ **425 Can't open data connection**
- ❖ **452 Error writing file**

# Electronic Mail

## Three components:

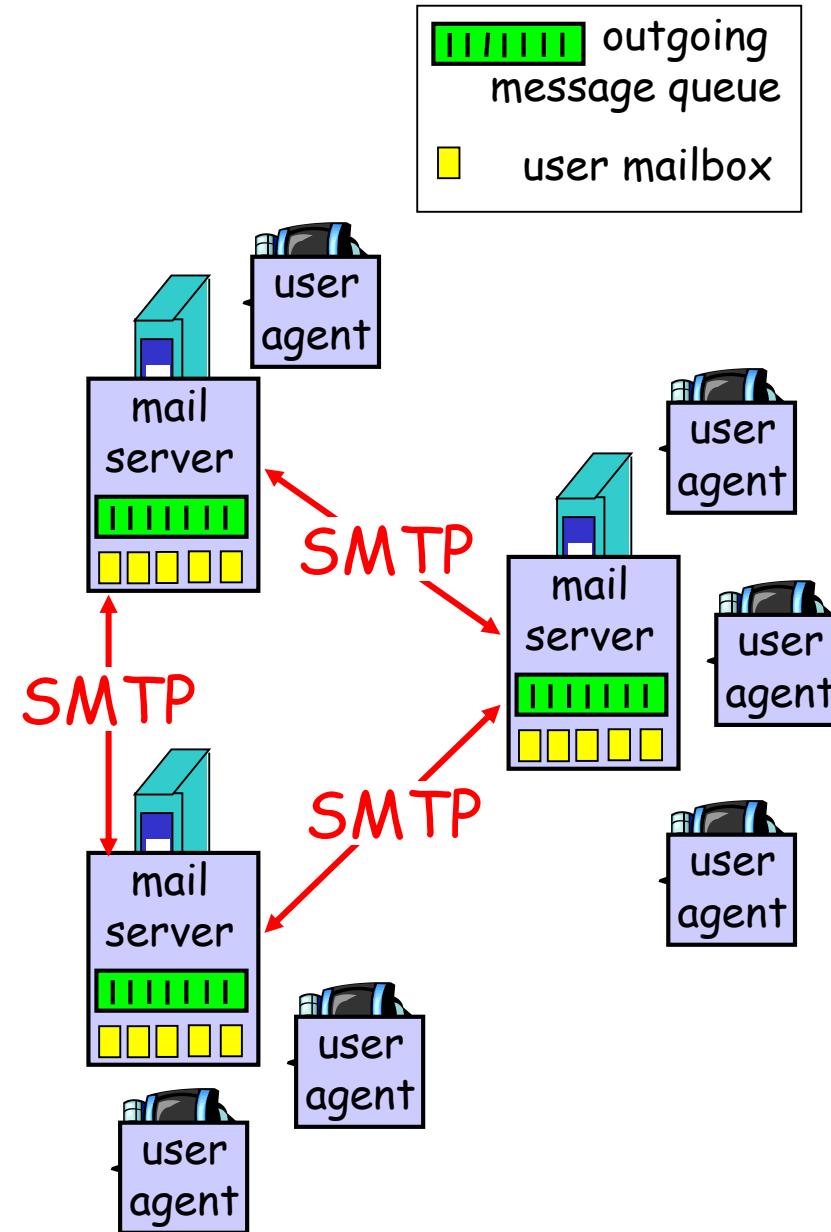
- 1. user agents, 2. mail servers
- 3. SMTP (simple mail transfer protocol)

### User Agent

- “mail reader”: editing, reading mail
- e.g., Outlook, Mozilla Thunderbird
- Out/incoming msgs stored on server

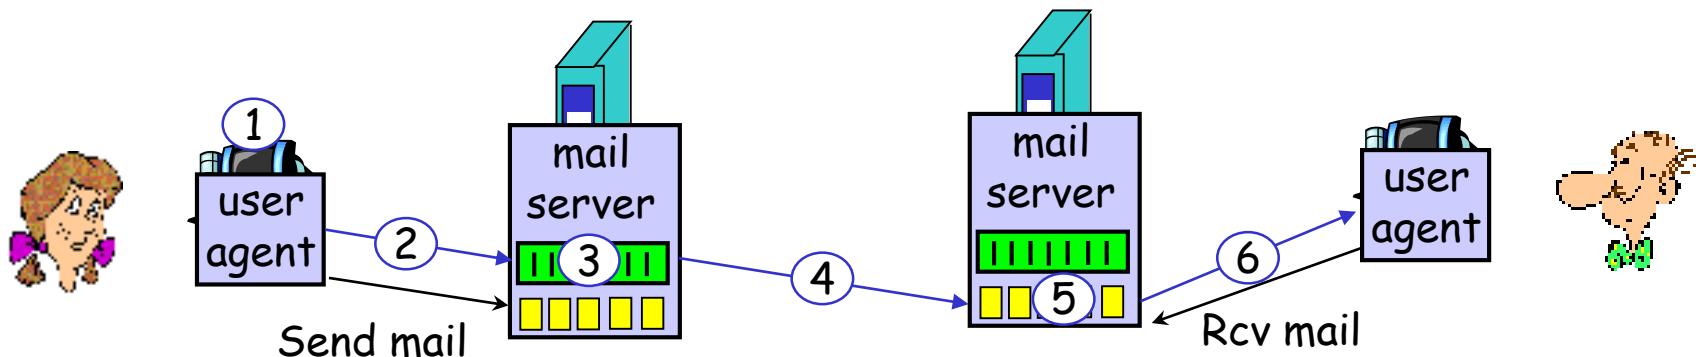
### Mail Servers

- **Mailbox**: incoming messages
- **message queue** outgoing msgs
- **SMTP protocol** between mail servers to send email messages
  - ❖ client: sending mail server
  - ❖ “server”: receiving mail server



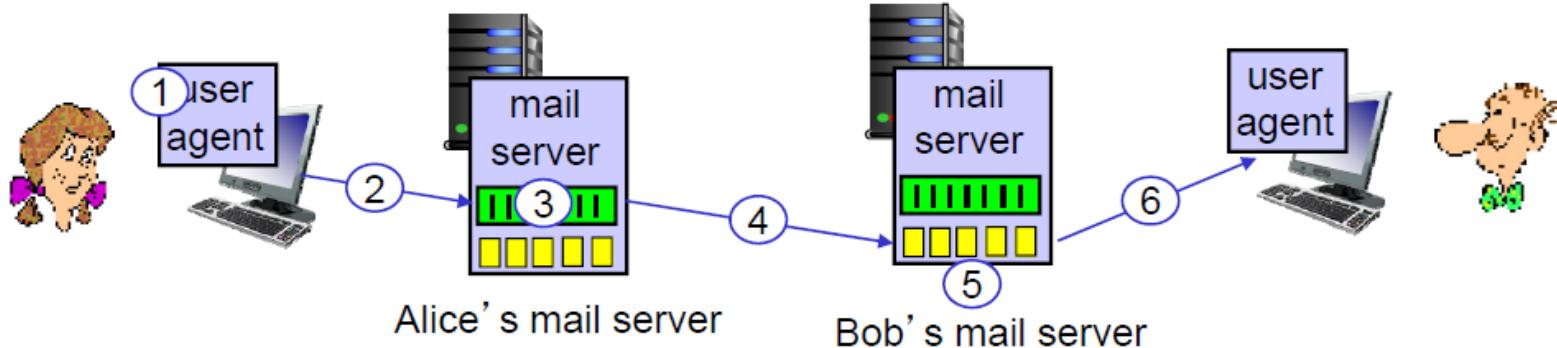
# Electronic Mail: SMTP [RFC 2821]

- uses TCP to reliably transfer email message from client to server, port 25
- direct transfer: sending server to receiving server
- three phases of transfer
  - ❖ 1. handshake, 2. transfer of messages, 3. closure
- SMTP uses persistent connections: sending mail server sends all its messages to the receiving mail server over one TCP connection
- Email Scenario:



# Scenario: Alice sends message to Bob

- 1) Alice uses UA to compose message “to” bob@someschool.edu
- 2) Alice’s UA sends message to her mail server; message placed in message queue
- 3) client side of SMTP opens TCP connection with Bob’s mail server
- 4) SMTP client sends Alice’s message over the TCP connection
- 5) Bob’s mail server places the message in Bob’s mailbox
- 6) Bob invokes his user agent to read message



## Try SMTP interaction for yourself:

- `telnet servername 25`
- see 220 reply from server
- enter HELO, MAIL FROM, RCPT TO, DATA, QUIT commands

above lets you send email without using email client  
(reader)

# SMTP: final words

- SMTP uses persistent connections
- SMTP requires message (header & body) to be in 7-bit ASCII
- SMTP server uses CRLF . CRLF to determine end of message

## *comparison with HTTP:*

- HTTP: pull
- SMTP: push
- both have ASCII command/response interaction, status codes
- HTTP: each object encapsulated in its own response message
- SMTP: multiple objects sent in multipart message

# Chapter 2: Application layer

- 2.1 Principles of network applications
- 2.2 Web and HTTP
- 2.3 FTP
- 2.4 Electronic Mail
  - ❖ SMTP, POP3, IMAP
- 2.5 DNS
- 2.6 P2P file sharing
- 2.7 Socket programming with TCP
- 2.8 Socket programming with UDP
- 2.9 Building a Web server

# DNS: Domain Name System

Internet identifiers for hosts, routers:

- ❖ IP address used for addressing datagrams
- ❖ "name", e.g., `www.yahoo.com`
  - used by humans

Q: map between IP addresses and name ?

Domain Name System:

- *distributed database* implemented in hierarchy of many *name servers*
- *application-layer protocol* host, routers, name servers to communicate to *resolve* names (address/name translation)
  - ❖ note: core Internet function, implemented as application-layer protocol
  - ❖ complexity at network's "edge"

# DNS

## DNS services

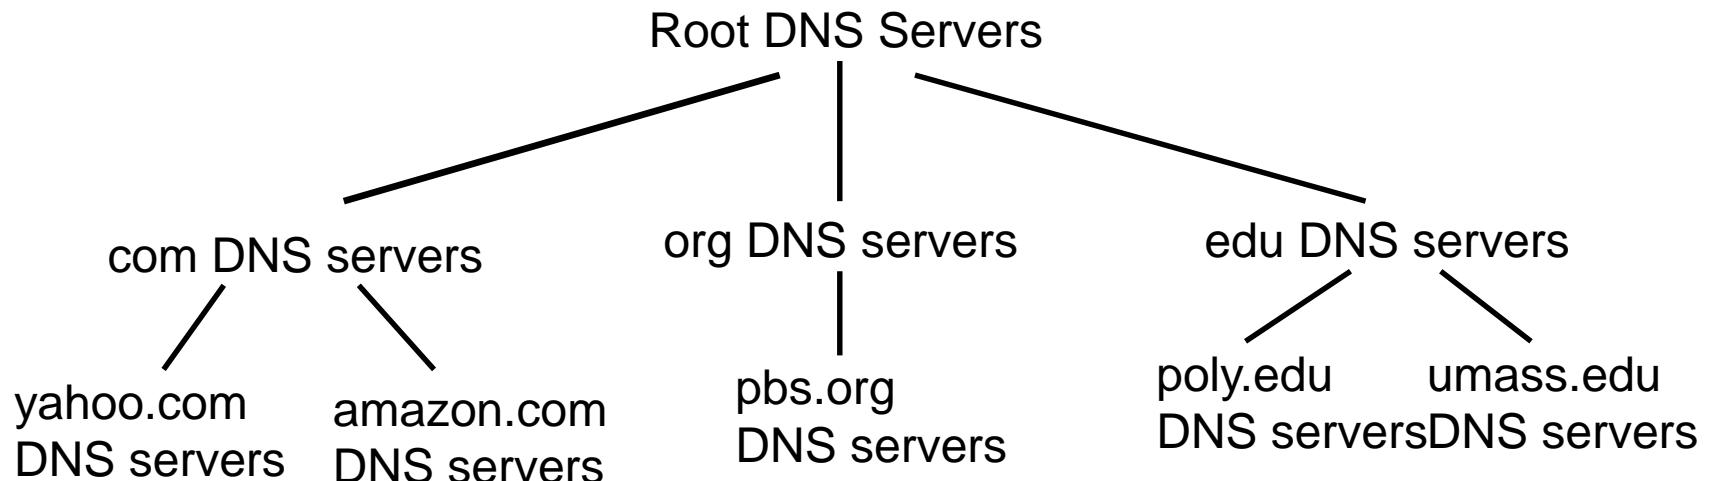
- hostname to IP address translation
- host aliasing
  - ❖ Canonical, alias names
- mail server aliasing
- load distribution
  - ❖ replicated Web servers: set of IP addresses for one canonical name

## Why not centralize DNS?

- single point of failure
- traffic volume
- distant centralized database = delays
- maintenance

*doesn't scale!*

# Distributed, Hierarchical Database

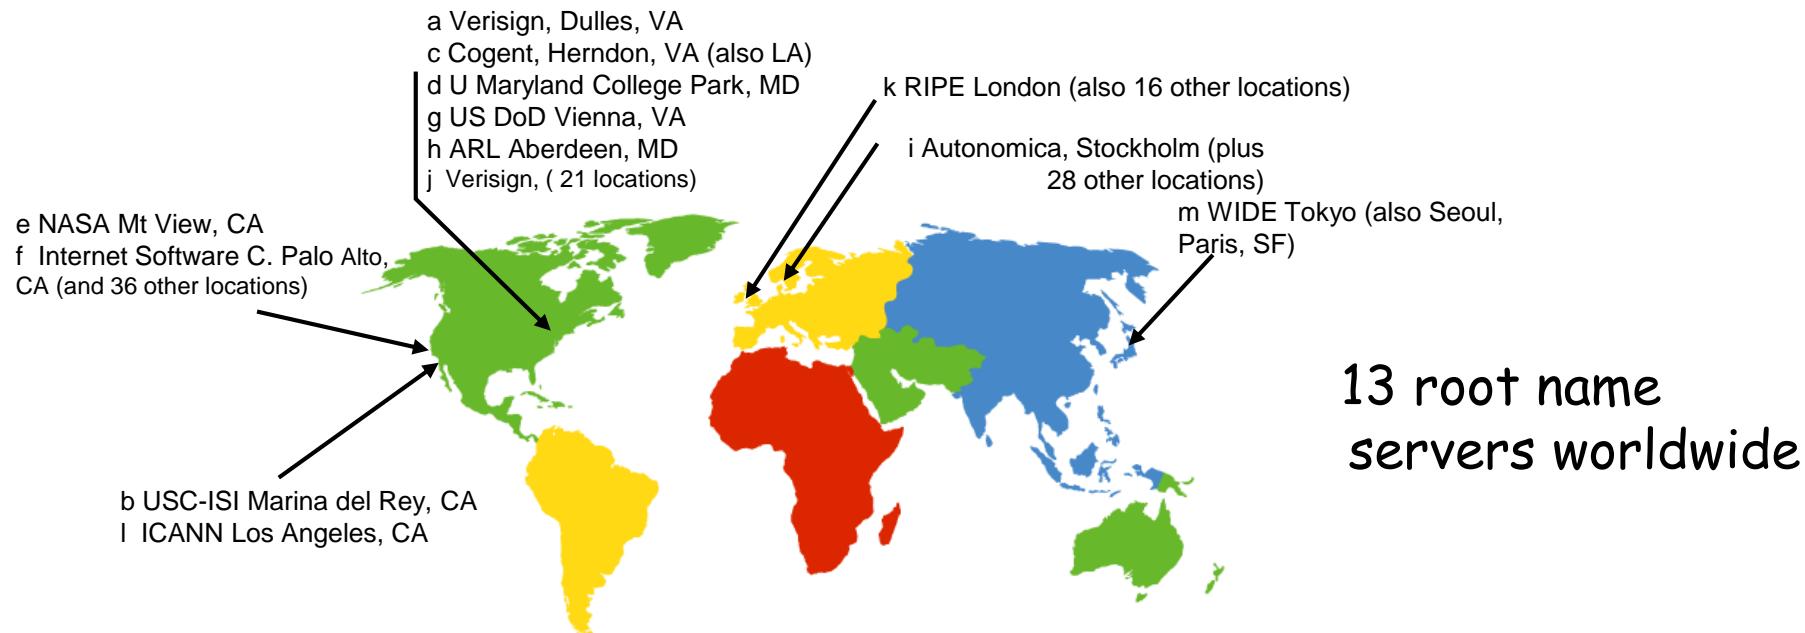


Client wants IP for www.amazon.com; 1<sup>st</sup> approx:

- client queries a root server to find com DNS server
- client queries com DNS server to get amazon.com DNS server
- client queries amazon.com DNS server to get IP address for www.amazon.com

# DNS: Root name servers

- contacted by local name server that can not resolve name
- root name server:
  - ❖ contacts authoritative name server if name mapping not known
  - ❖ gets mapping
  - ❖ returns mapping to local name server



# TLD and Authoritative Servers

## □ I. Top-level domain (TLD) servers:

- ❖ responsible for com, org, net, edu, etc, and all top-level country domains uk, fr, ca, jp.
- ❖ Network Solutions maintains servers for com TLD
- ❖ Educause for edu TLD

## □ II. Authoritative DNS servers:

- ❖ organization's DNS servers, providing authoritative hostname to IP mappings for organization's servers (e.g., Web, mail).
- ❖ can be maintained by organization or service provider

### III. Local Name Server

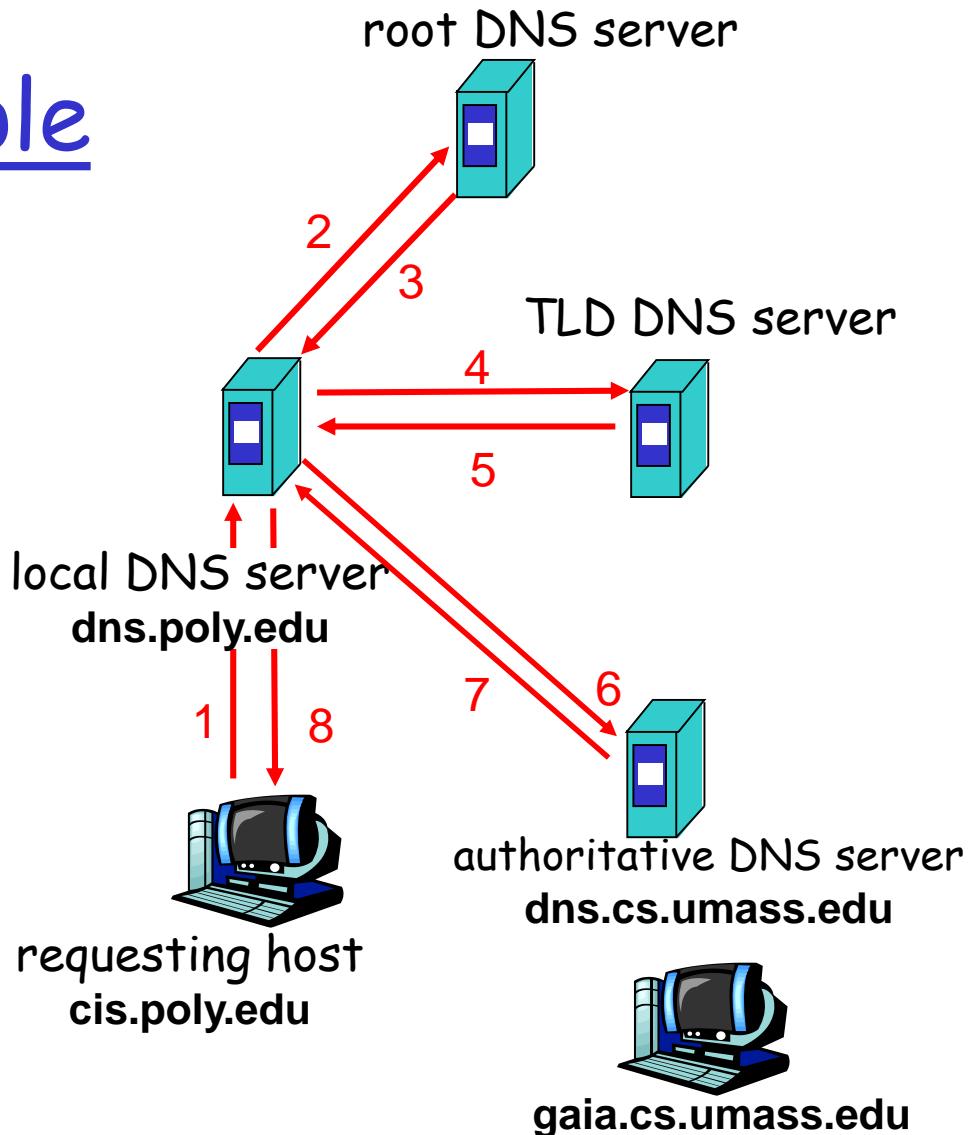
- does not strictly belong to hierarchy
- each ISP (residential ISP, company, university) has one.
  - ❖ also called "default name server"
- when host makes DNS query, query is sent to its local DNS server
  - ❖ acts as proxy, forwards query into hierarchy

# DNS name resolution example

- Host at cis.poly.edu wants IP address for gaia.cs.umass.edu

## A. iterative query:

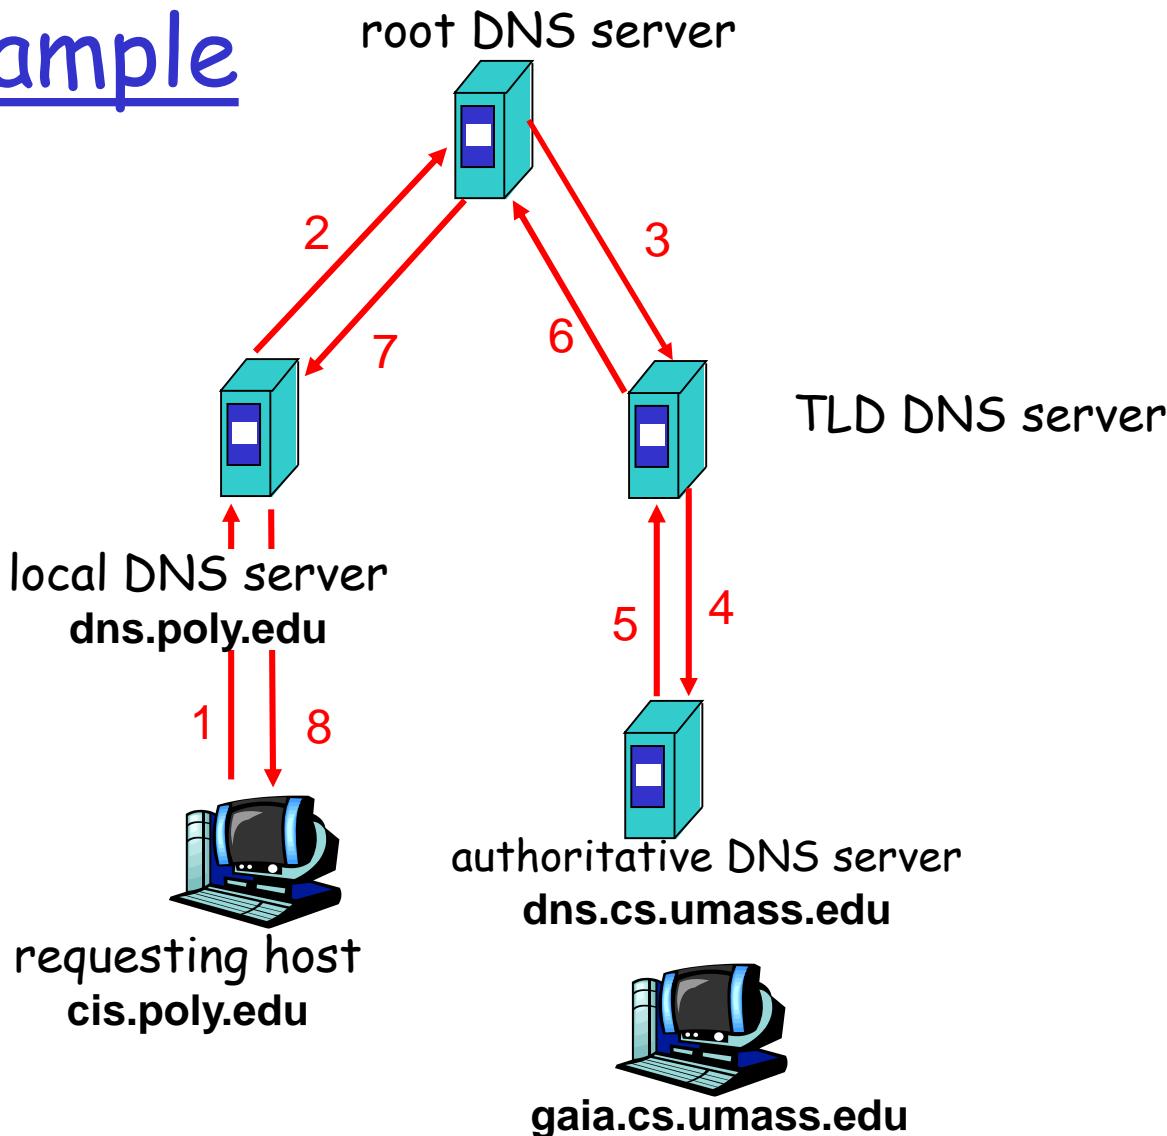
- contacted server replies with name of server to contact
- “I don’t know this name, but ask this server”



# DNS name resolution example

## B. recursive query:

- puts burden of name resolution on contacted name server
- heavy load?



# Chapter 2: Application layer

- 2.1 Principles of network applications
  - ❖ app architectures
  - ❖ app requirements
- 2.2 Web and HTTP
- 2.4 Electronic Mail
  - ❖ SMTP, POP3, IMAP
- 2.5 DNS
- 2.6 P2P file sharing
- 2.7 Socket programming with TCP
- 2.8 Socket programming with UDP
- 2.9 Building a Web server

# P2P file sharing

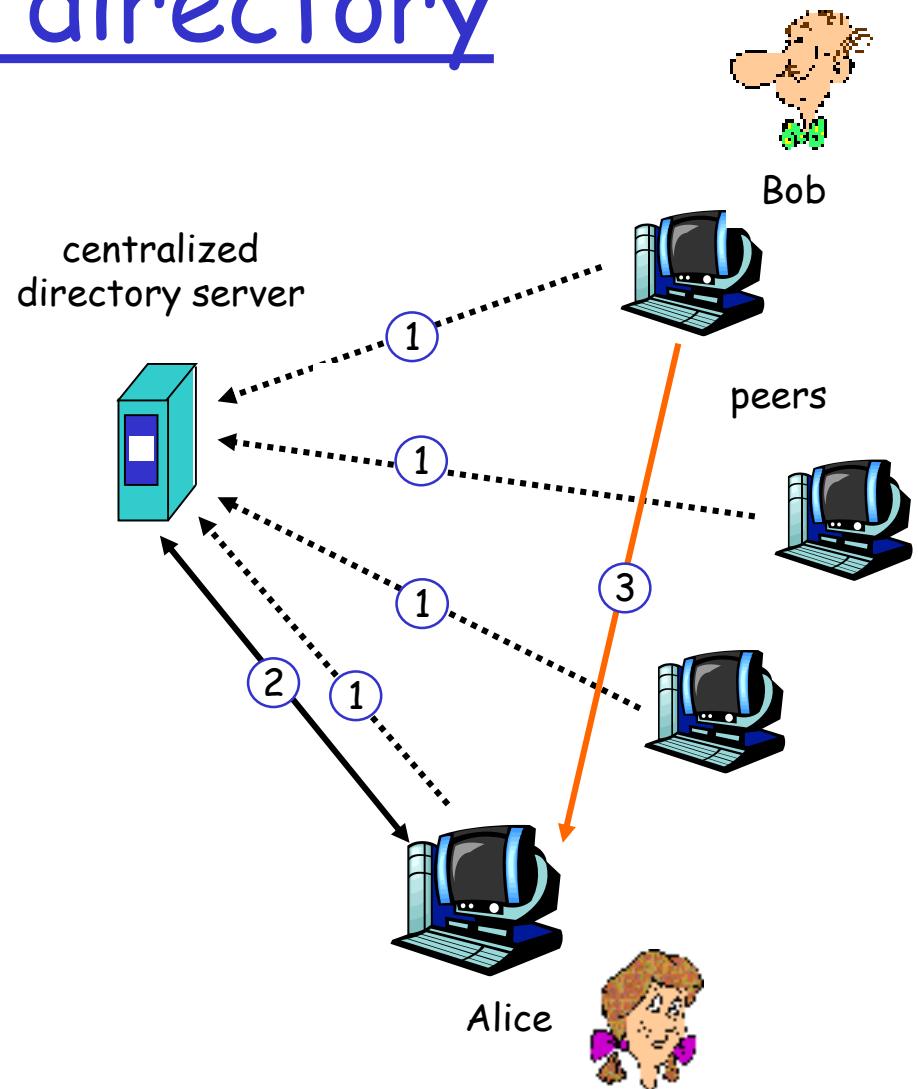
## Example

- Alice runs P2P client application on her notebook computer
  - intermittently connects to Internet; gets new IP address for each connection
  - asks for "Hey Jude"
  - application displays other peers that have copy of Hey Jude.
- Alice chooses one of the peers, Bob.
  - file is copied from Bob's PC to Alice's notebook: HTTP
  - while Alice downloads, other users uploading from Alice.
  - Alice's peer is both a Web client and a transient Web server.
- All peers are servers = highly scalable!

# P2P: centralized directory

original “Napster” design

- 1) when peer connects, it informs central server:
  - ❖ IP address
  - ❖ content
- 2) Alice queries for “Hey Jude”
- 3) Alice requests file from Bob



# P2P: problems with centralized directory

- single point of failure
- performance bottleneck
- copyright infringement:  
“target” of lawsuit is  
obvious

file transfer is decentralized, but locating content is highly centralized

Advantages vs. disadvantages  
Search time and overhead?